

BT5153

Applied Machine Learning for Business Analytics

ZHAO Rui
`diszr@nus.edu.sg`

<https://bt5153msba.github.io/>

Course Lecturer

- Quantitative Researcher in a hedge fund
- Adjunct Faculty at NUS, teaching BT5153 and BT4012
- Research interests are within machine learning and its applications on time series data and text data.
- [Google Scholar](#)
- [Personal Homepage](#)



Zhao Rui
diszr@nus.edu.sg

Course TAs



Sanjay Saha
sanjaysaha@u.nus.edu



Xutao Wang
wangxutao@u.nus.edu

Agenda

1. Course Overview
2. What is Machine Learning?
3. Machine Learning is everywhere
4. Overview of machine learning
5. Group Project

What BT5153 covers

Goals of this Course

Learn and improve upon the applications of machine learning

- Understand conceptually the mechanism of machine learning algorithms
- Implement the whole data science pipeline
- Select appropriate machine learning tools/techniques for business applications.

Course Background and Overview

- After DSC5106 Foundation in Data Analytics I
- Together with BT5151 Foundation in Data Analytics II
- Basic Machine Learning/Data Mining models have been covered in these two above modules
- In BT5153:
 - “**Advanced**” topics
 - **Hands-on** Experiences
 - In each lecture, roughly 75% Slides and 25% IPython notebooks.
 - More **Practical** Assignments/Exams

Models

- Representation Learning
 - Autoencoder
 - Word Embeddings
 - BERT
- Bayesian Learning
 - Naive Bayes
 - Bayesian Linear Regression
- Deep Learning
 - Neural Networks
 - Convolutional Neural Network
 - Transfer Learning
- Explainable Machine Learning

Applications

- Spam Detection
- Document Classification
- Recommendation
- Image Categorization
- Sentiment Analysis
- Image/Text Generation
- Question Answering Tasks
- Name Entity Recognition
- Part-of-Speech Tagging
- Etc

Hands-on Experience

- **Understanding domain, prior knowledge**
- **Date integration, selection, clearing, pre-processing, etc**
- Learning models (little math, more intuitive ideas)
- Compare models
- **Model Interpretability**
- Consolidating and deploying discovered knowledge
- **Apply discovered knowledge to practical problems**
- Python programming is not the teaching focus.

Course Assessment

- In-class Quizzes (10%)
- Individual Assignments (40%):
 - Two weekly individual assignments (each 10%)
 - One mini-kaggle project (20%)
- Group Project (50%)
 - Project Proposal (10%)
 - Final Presentation (20%)
 - Final Report (20%)

Course Schedule

Date	Topic	Content	Assignment
Fri 01/15	Introduction to Machine Learning	LINK	N.A.
Fri 01/22	Machine Learning Practices	LINK	N.A.
Fri 01/29	Bayesian Models	LINK	Assignment I Out
Fri 02/05	Neural Networks and Deep Learning	LINK	Form your team
Fri 02/12	Chinese New Year	LINK	Group Project Proposal Due
Fri 02/19	Deep Learning Practices	LINK	Assignment II Out
Fri 02/26	Recess Week	N.A.	N.A.
Fri 03/05	AutoEncoder	LINK	Kaggle Starts
Fri 03/12	Convolutional Neural Networks	LINK	N.A.
Fri 03/19	Explainable Machine Learning	LINK	N.A.
Fri 03/26	Frontiers in NLP: I	LINK	Kaggle Competition Due
Fri 04/02	Good Friday	LINK	Kaggle Report Due
Fri 04/09	Frontiers in NLP: II	N.A.	N.A.
Fri 04/16	Responsible Machine Learning and Course Summary	LINK	Presentation Recording Due
Fri 04/23	Reading Week	N.A.	Final Report Due

How to get Slides & Notebooks

Date	Topic	Content	Assignment
Fri 01/17	Introduction to Machine Learning	LINK	N.A.
Fri 01/24	Machine Learning Practice	LINK	N.A.
Fri 01/31	Explainability-Accuracy Tradeoff	LINK	Form your team
- -			

Week 1

In-class Material

1. [Slides](#)
2. [Notebook](#)

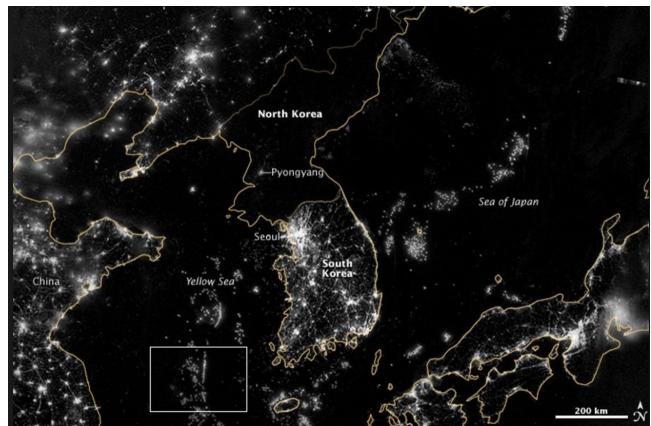
Extra Reading

1. [Google Research: Looking Back at 2019, and Forward to 2020 and Beyond](#)
2. [Hopes from AI experts for 2020](#)
3. [AI & ML based App Ideas](#)
4. [AI can show us the ravages of climate change](#)
5. [Ted Talk: Image Recognition by Prof. Li Feifei](#)
6. [Truly Master Scikit-Learn](#)

The content for each class will be updated a week ahead of the schedule.



Face Recognition



GDP Prediction



Hedge fund use ML for trading



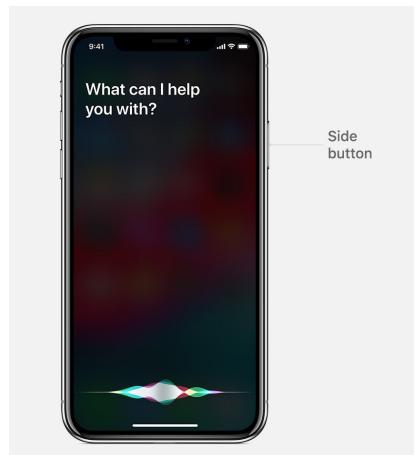
Special Effects in Tiktok



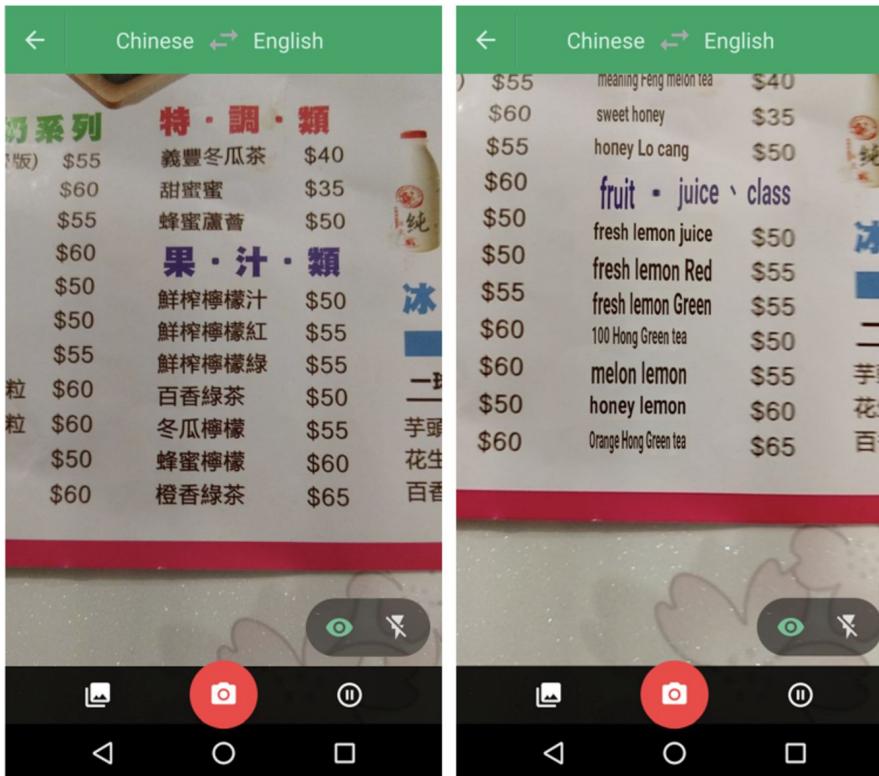
Amazon Recommendation



AlphaGo



Siri



Machine Translation



Self-driving Car

What is Machine Learning



Mat Velloso

@matvelloso

Follow



Difference between machine learning
and AI:

If it is written in Python, it's probably
machine learning

If it is written in PowerPoint, it's
probably AI

5:25 PM - 22 Nov 2018

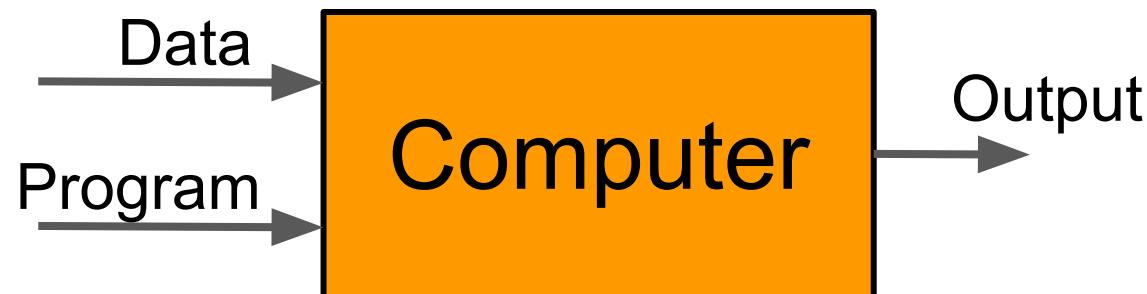
8,541 Retweets **23,778** Likes



Python Programming

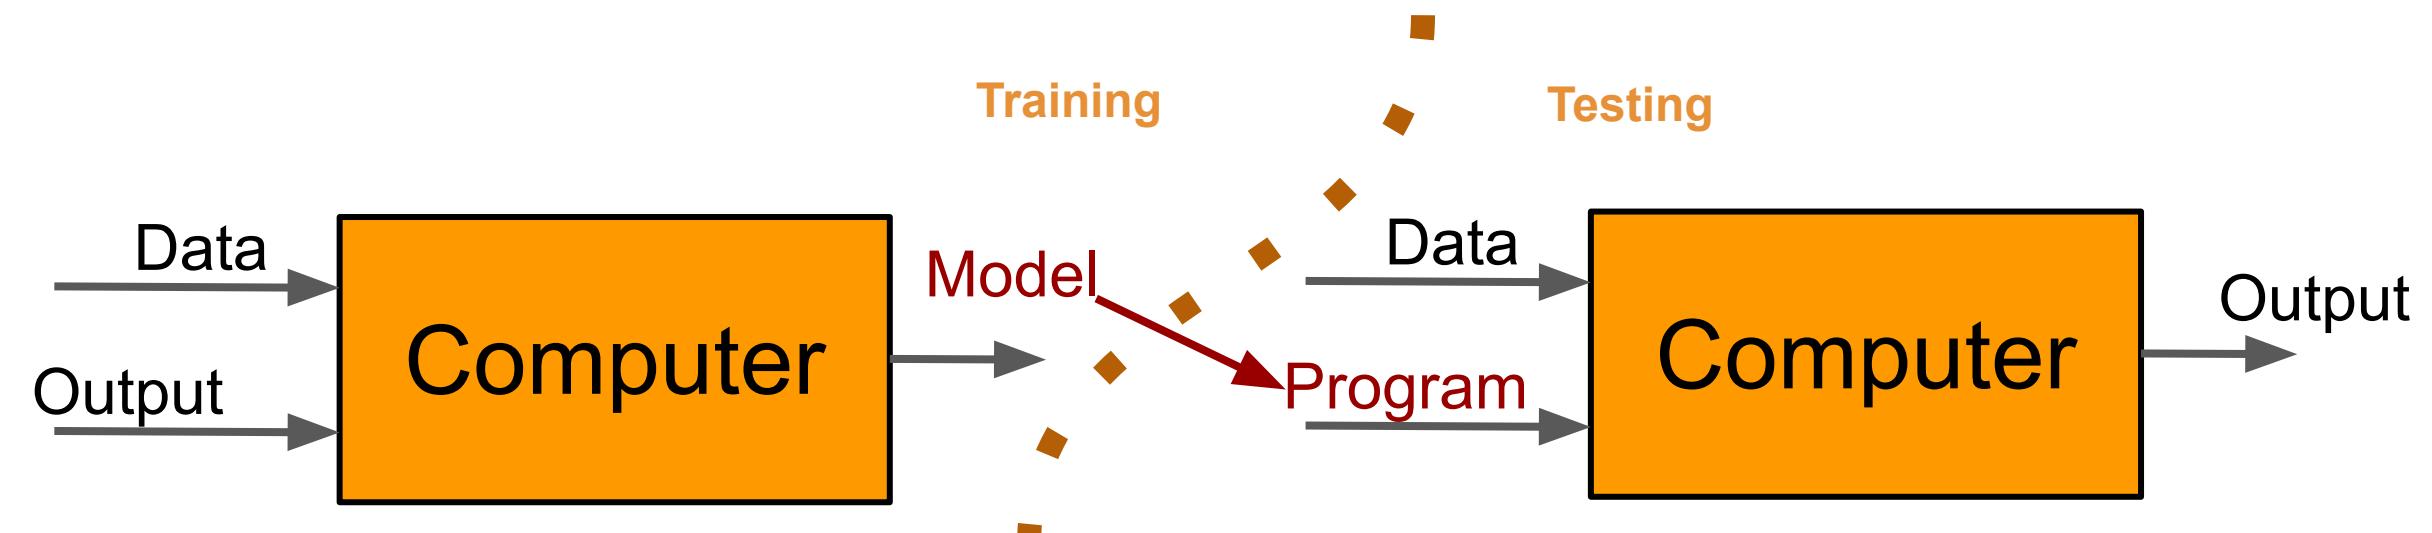
```
In [1]: a = 3  
       b = 1  
       q = 3*a + 2*b  
       print('result is {}'.format(a + b))
```

result is 4



Machine Learning

```
: from sklearn.neighbors import KNeighborsClassifier  
from sklearn.metrics import accuracy_score  
#create an object of KNN  
neigh = KNeighborsClassifier(n_neighbors=3)  
#train the algorithm on training data and predict using the testing data  
pred = neigh.fit(data_train, target_train).predict(data_test)
```



Definition of Machine Learning

“A computer program is said to learn from **experience E** with respect to some class of **tasks T** and performance **measure P**, if its performance at tasks in **T**, as measured by **P**, improves with experience **E**”



Tom Mitchell

T, **P**, **E** are three basic elements to define a complete machine learning tasks

AlphaGo

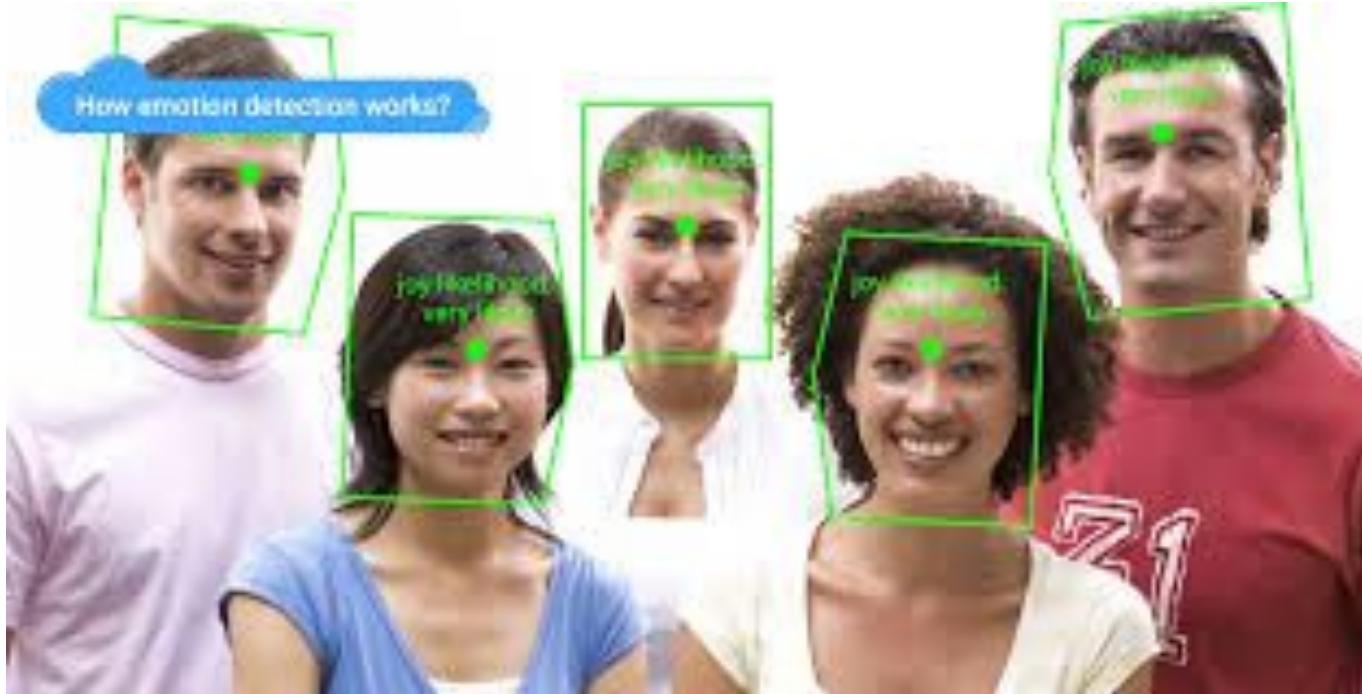


T: Play Go Games

P: Win rates of all matches

E: Match Experiences with many go players or itself

Face Recognition



- T**: Identify or verify human faces
- P**: Accuracy that human faces are detected
- E**: Dataset of labelled human faces

Machine Translation



- T**: Translate source language into target language
- P**: Accuracy that the language have been translated
- E**: Corpus of source-target language pairs

More about E

- For machine learning algorithms, E is **data**.

- Data types:

- Unstructured vs Structured
- Raw vs Processed



Computers Processable and Understandable

The very import step for Machine Learning Project is that how to **preprocess these unstructured/raw data**.

Structured

- Structured: Table (Matrix) or Tensor

Player	Height (inches)	Weight (pounds)	Position
Player 1	76	225	C
Player 2	75	195	PG
Player 3	72	180	SF
Player 4	82	231	PF

Features → Height (inches), Weight (pounds)

Feature Values → 76, 225, 75, 195, 72, 180, 82, 231

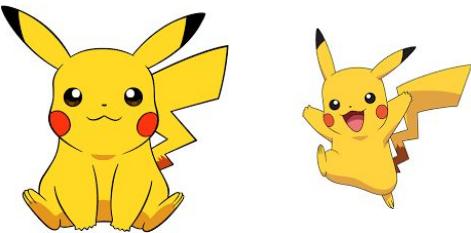
Data Sample → Player 4

Labels → Position

Unstructured

- The original data can not be stored in an “table”
- More abstract, more fuzzy, and more high-dimensionality

Images



Audio



Video



Text

Content

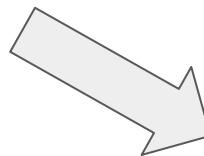
This module provides students a deep overview of various advanced machine learning techniques applied to business analytics tasks. The focus of this course will be the key and intuitive idea behind machine learning models and hands-on examples instead of theoretical analysis. The tentative topics include machine learning pipeline, unsupervised learning, structure learning, Bayesian learning, deep learning and generative models. The programming languages used will be Python.

Environment around agent



Raw Data

- Unstructured Data such as images and text
- Some structured data like categorical data



Labels

Position

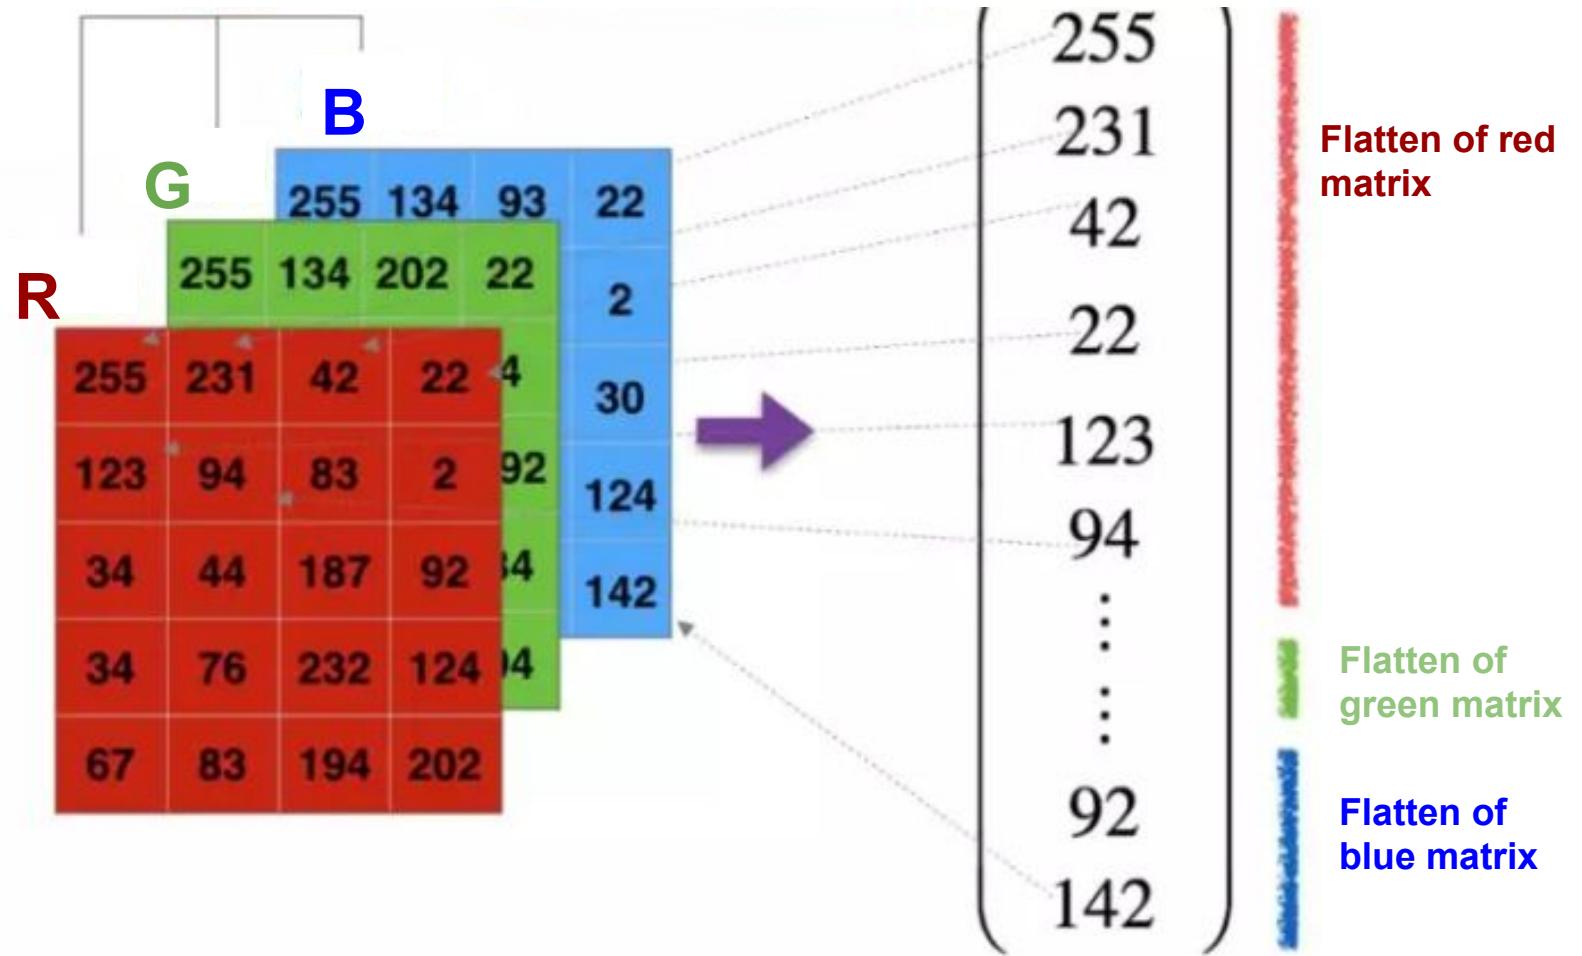
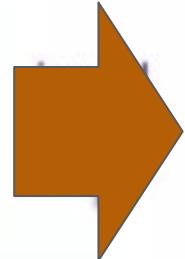
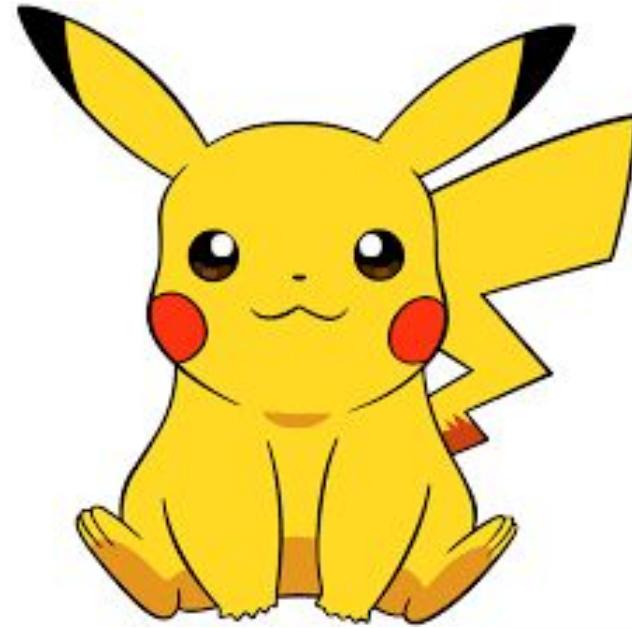
C

PG

SF

PF

Processed Data (from Raw)



Processed Data (from Raw)

the dog is on the table

0	0	1	1	0	1	1	1
are	cat	dog	is	now	on	table	the

Processed Data (from Raw)

Position	Label	Position	Label
C	0	C	1 0 0 0
PG	1	PG	0 1 0 0
SF	2	SF	0 0 1 0
PG	1	PG	0 0 1 0
PF	3	PF	0 0 0 1

Which one is better?

Terms

- Artificial Intelligence: **Intelligence** exhibited by machines to mimic a human mind
- Machine Learning: Computers being able to learn without hand-coding each step
- Deep Learning: **Multi-layered** algorithms for learning from data
- Data Science: Methods, processes, and systems to extract **insights** from data
- Analytics: Discovery of meaningful patterns in data

What is what

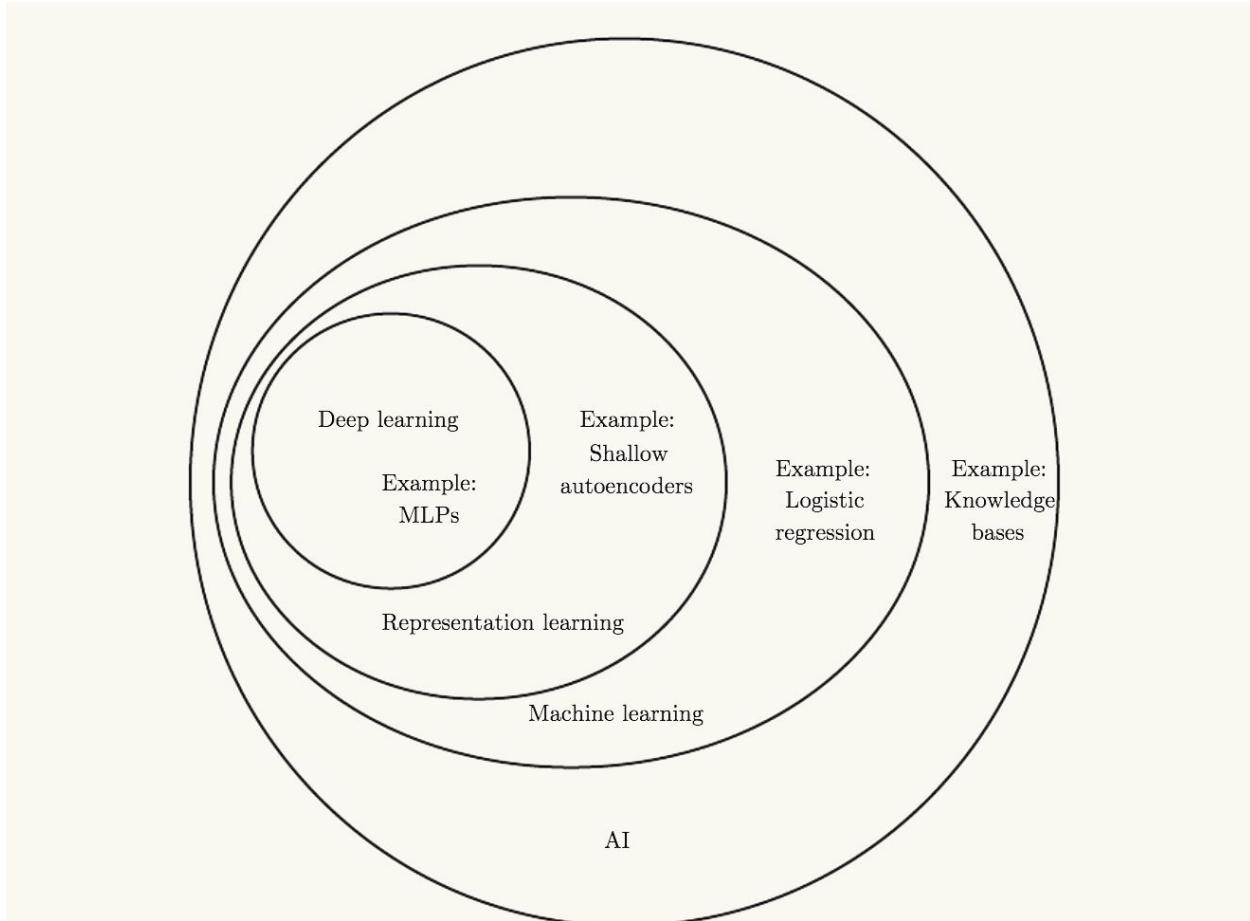


Figure 1.4: A Venn diagram showing how deep learning is a kind of representation learning, which is in turn a kind of machine learning, which is used for many but not all approaches to AI. Each section of the Venn diagram includes an example of an AI technology.

Goodfellow, Ian, et al. *Deep learning*. Vol. 1. Cambridge: MIT press, 2016.

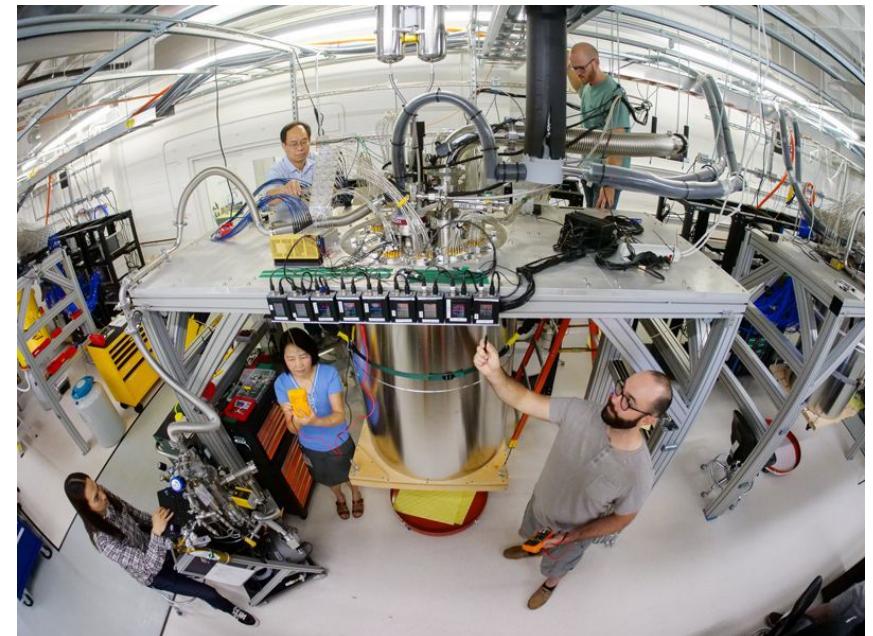
Machine Learning is
everywhere

Why Study Machine Learning?

- Reduce Human Efforts:
 - Allow the computers learn automatically
 - Writing rule-based program is too **challenging**
 - Let the data **SPEAK**
- Machine learning brings better career opportunity
 - The trend will be **MLaaS** (Machine learning as a service)
 - Machine learning engineer, Data scientist, Data Product Manager, Cloud Engineer, etc

Why Machine Learning is Powerful?

- Recent progress in algorithms and theory
- Big data era
 - Flow of online data and mobile data
 - 5G is developing
 - Cloud computing
- Computational power is available
 - TPU, GPU, **Quantum Computing**,

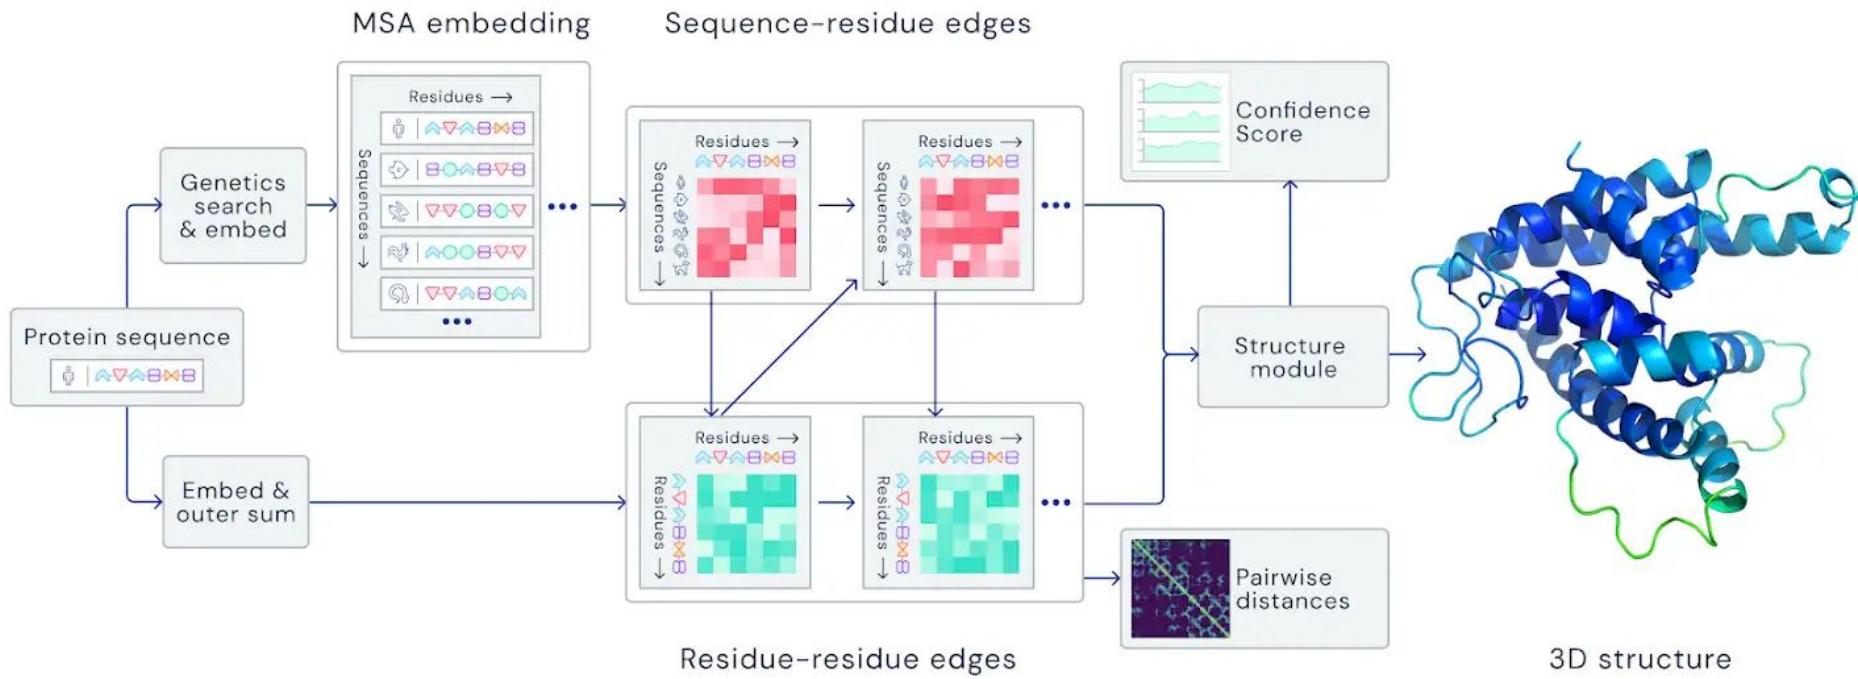


Google Quantum Computer

Three Niches for Machine Learning

- **Data mining**
 - Use historical data to improve decisions.
- **Software applications that are hard to be programmed by hand**
 - Speech Recognition
 - Autonomous Driving
 - Etc
- **User modeling**
 - Recommendation System
 - Micro-credit Loan
 - Etc

AlphaFold



source: <https://deepmind.com/blog/article/alphafold-a-solution-to-a-50-year-old-grand-challenge-in-biology>

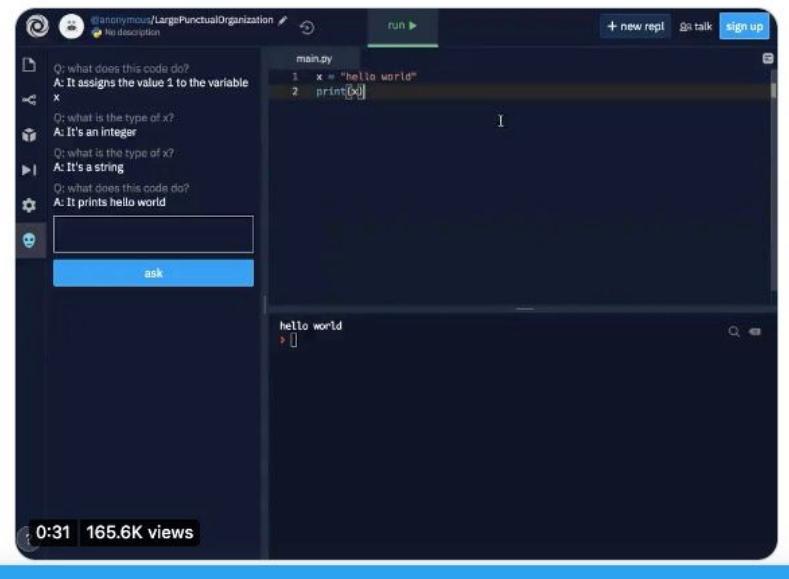
GPT-3: Generate Text

Reading code is hard! Don't you wish you could just ask the code what it does? To describe its functions, its types.

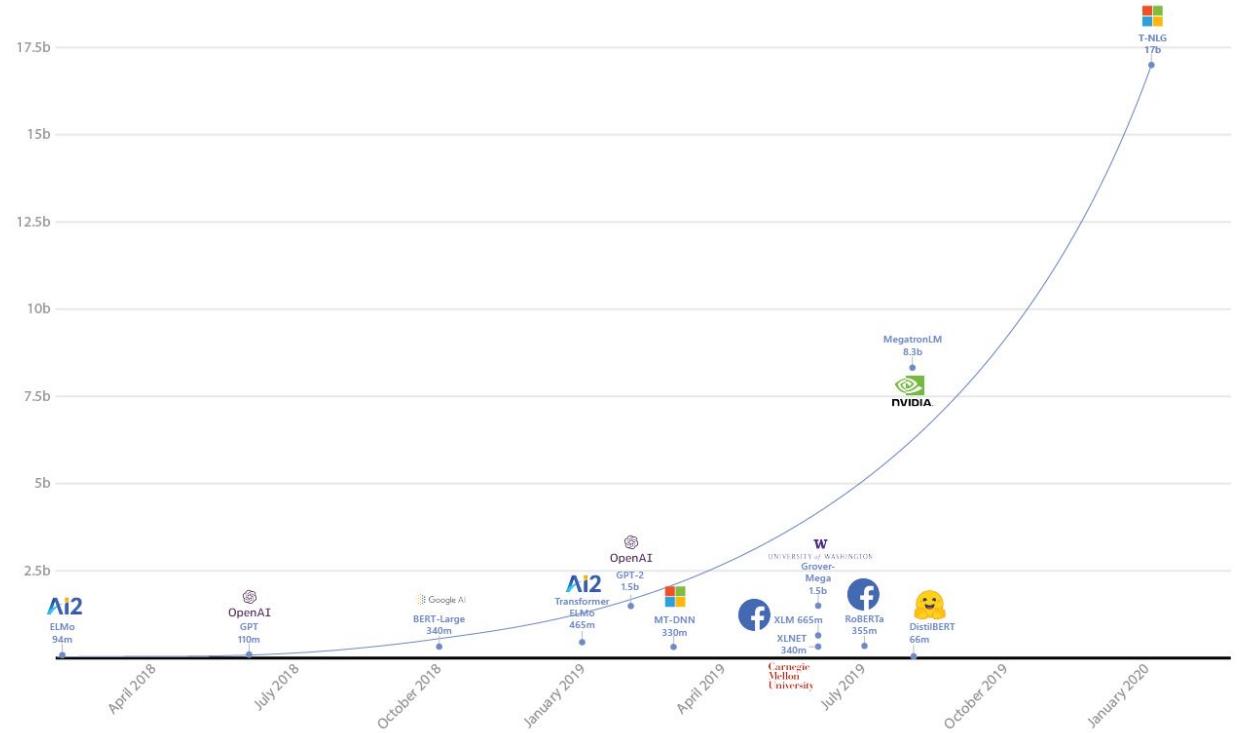
And maybe... how can it be improved?

Introducing: [@Replit](#) code oracle 

It's crazy, just got access to [@OpenAI](#) API and I already have a working product!



<https://twitter.com/i/status/1285789362647478272>

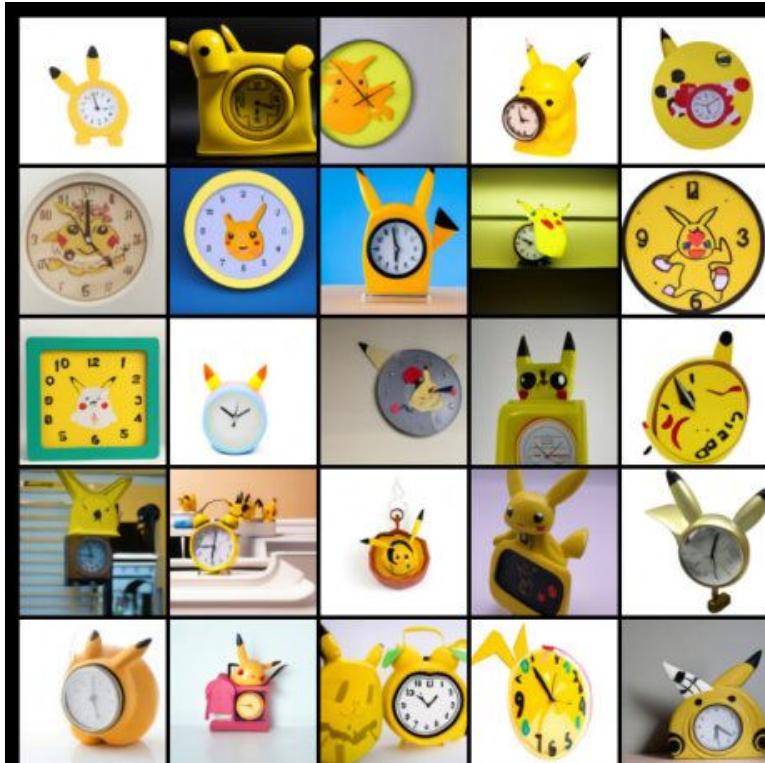


<https://www.microsoft.com/en-us/research/blog/turing-nlg-a-17-billion-parameter-language-model-by-microsoft/>

DALL-E: Creating Image from Text

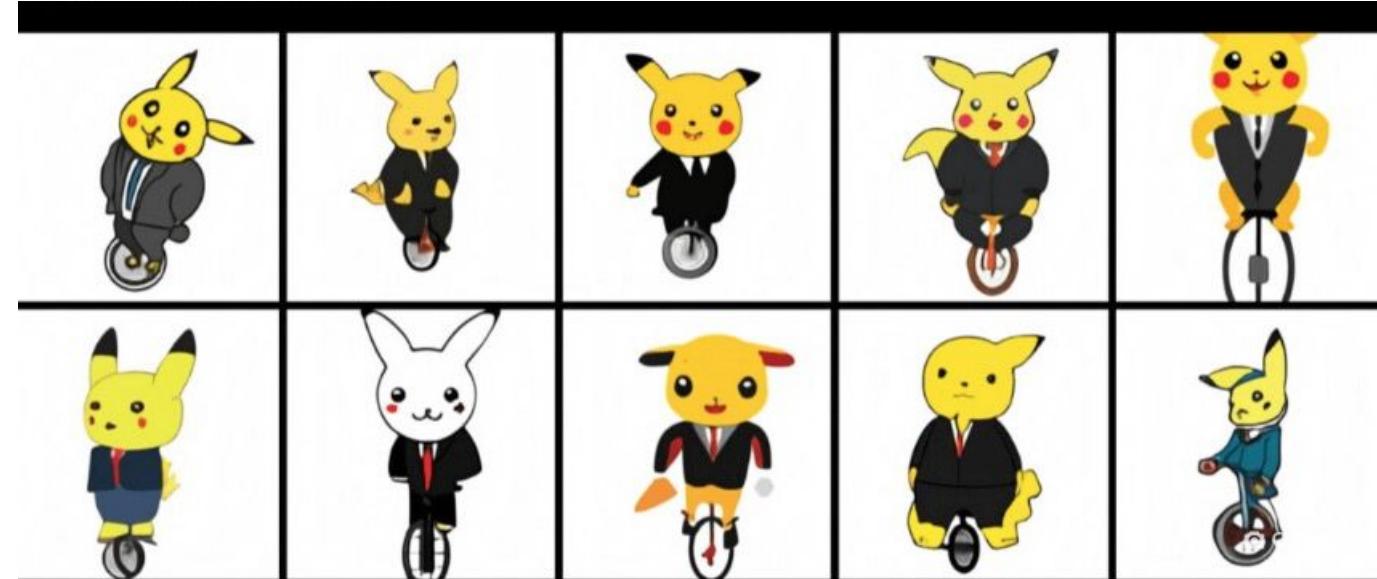
Input
Text

A clock in the style of a pikachu



Generated
Images

A pikachu in a suit riding a unicycle



source: <https://openai.com/blog/dall-e/>

Overview of Machine Learning Concepts

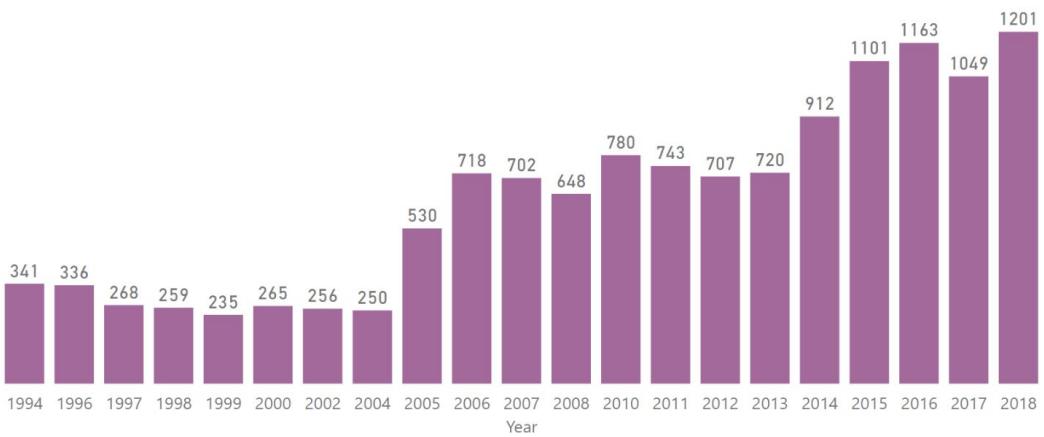
Basic Paradigm

- Define the **T**asks (what should be learned)
- Find the training **E**xperiences(datasets)
- Quantify the **P**erformance via a measure
- Choose a machine learning model to complete the **T** and improve the **P** via the **E**

Machine Learning in a nutshell

- Ten of thousands of machine learning algorithms

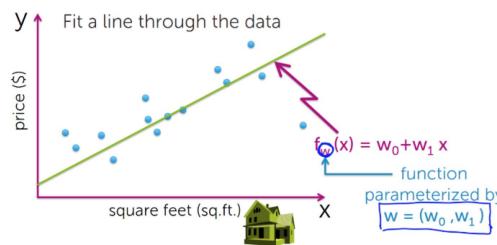
Publications per year



- Hundreds New per month

- Each ML algorithm can be decomposed into:
 - **Representation**
 - **Evaluation**
 - **Optimization**

<https://www.microsoft.com/en-us/research/project/academic/articles/aaai-conference-analytics/> AAAI Conference



Representation

- Decision Trees
- Support Vector Machine
- Set of rules
- Instances-based Learning (K Nearest Neighbor)
- Probabilistic Graph Models (Naive Bayes and Hidden Markov Models)
- Neural Networks and Deep Learning
- Ensemble
- Others

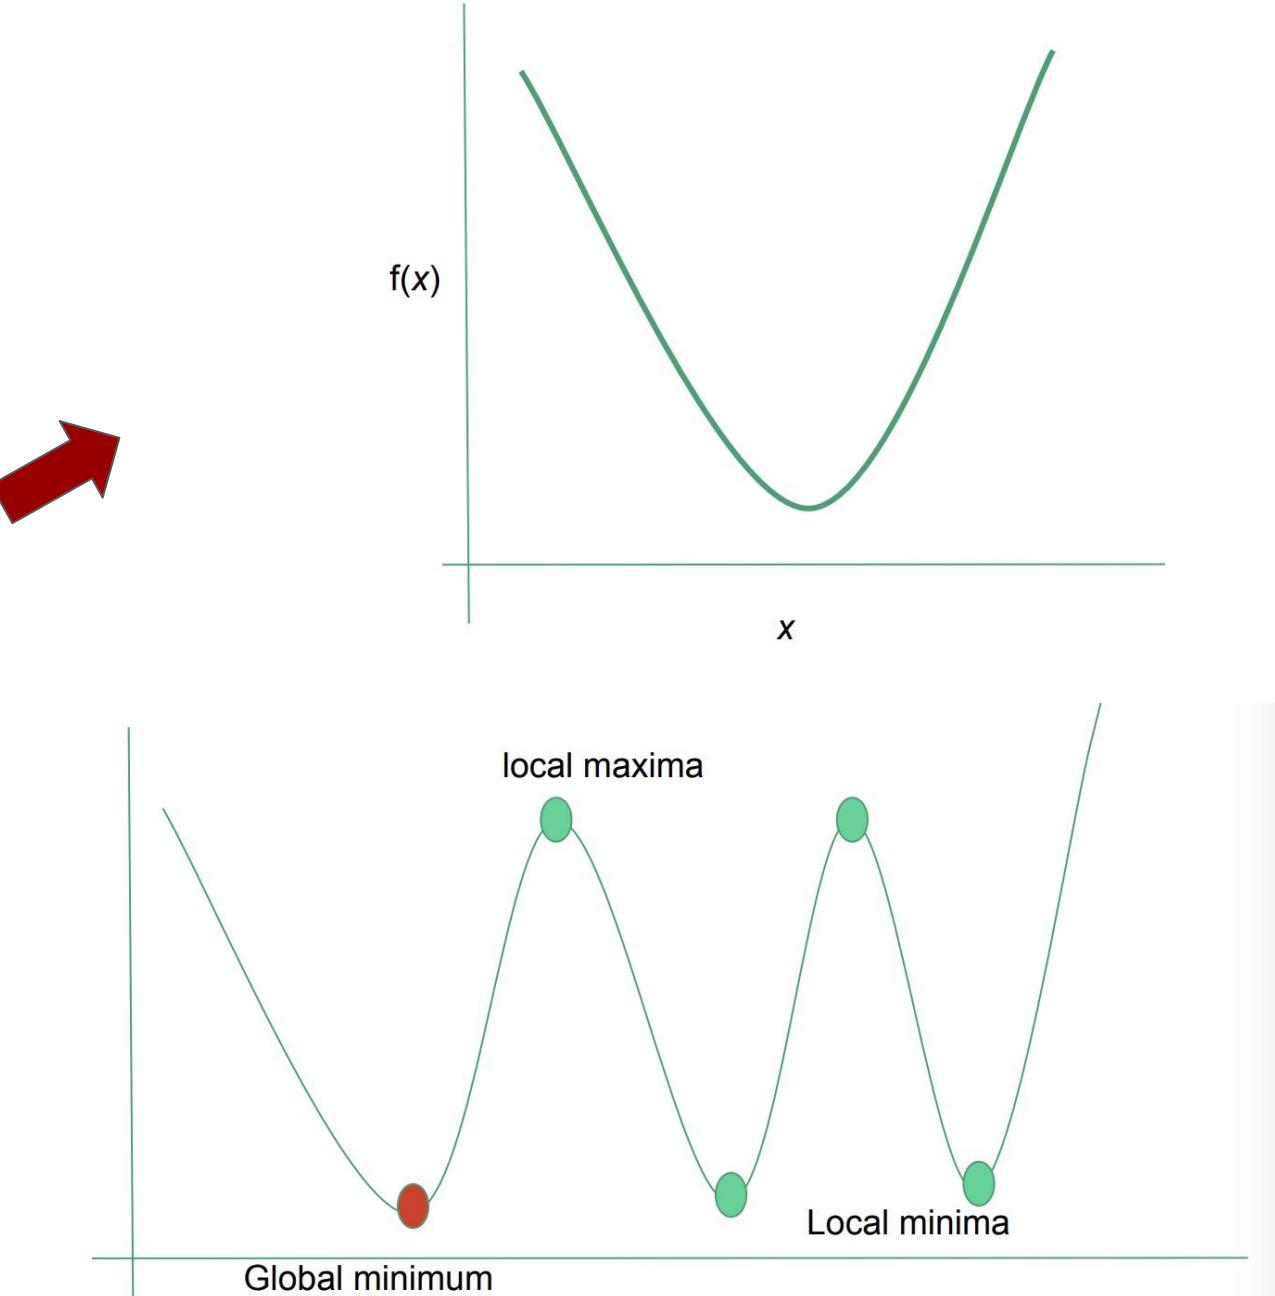
ATTENTION: this is not referred to the one in “representation learning”

Evaluation

- Confusion matrix
- Accuracy
- Precision and recall
- Mean Squared Error
- Likelihood
- Posterior probability
- Margin
- Entropy
- K-L Divergence
- Etc

Optimization

- Combinatorial Optimization
 - Grid search
- Convex Optimization
 - Least Squares
 - Linear Programming (with constraints)
 - Semidefinite Programming
 - Etc
- Non-convex Optimization
 - Gradient descent algorithm
 - Bayesian Optimization
 - Etc



Types of Machine Learning Models

- Supervised Learning
 - Training data contain the desired outputs (labels)
- Unsupervised Learning
 - Training data do not contain labels
- Reinforcement Learning
 - Rewards from **sequence actions**
- Semi-supervised Learning
 - Training data include a few labels

Supervised Learning

- Given (training data x , training label y), predict (new data, ?)

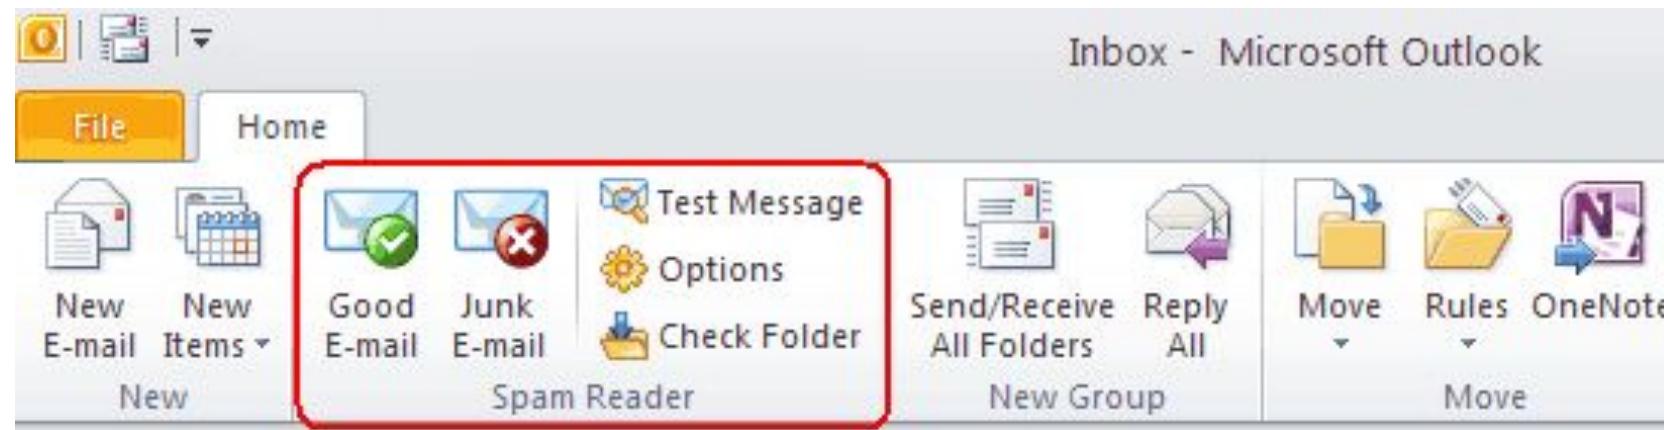
$$h(x) \approx f(x)$$



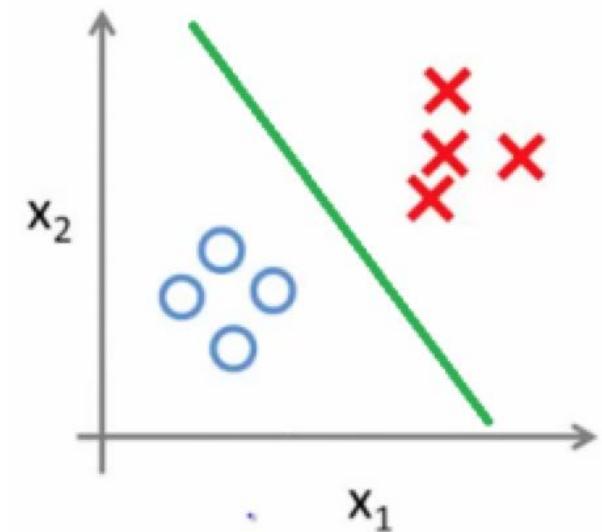
- Classification: **discrete** label
 - Binary Classification: label y in $\{0, 1\}$
 - Multi-class Classification: label y in $\{0, 1, 2, 3, \dots, k-1\}$
- Regression: **continuous** label
 - y is real-valued space

Binary Classification

- Junk Email Filter

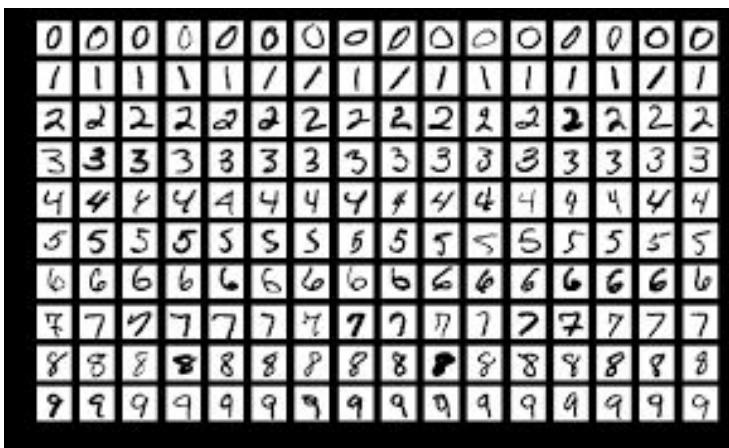


- Two classes
 - 1 Normal emails
 - 0 Spam emails

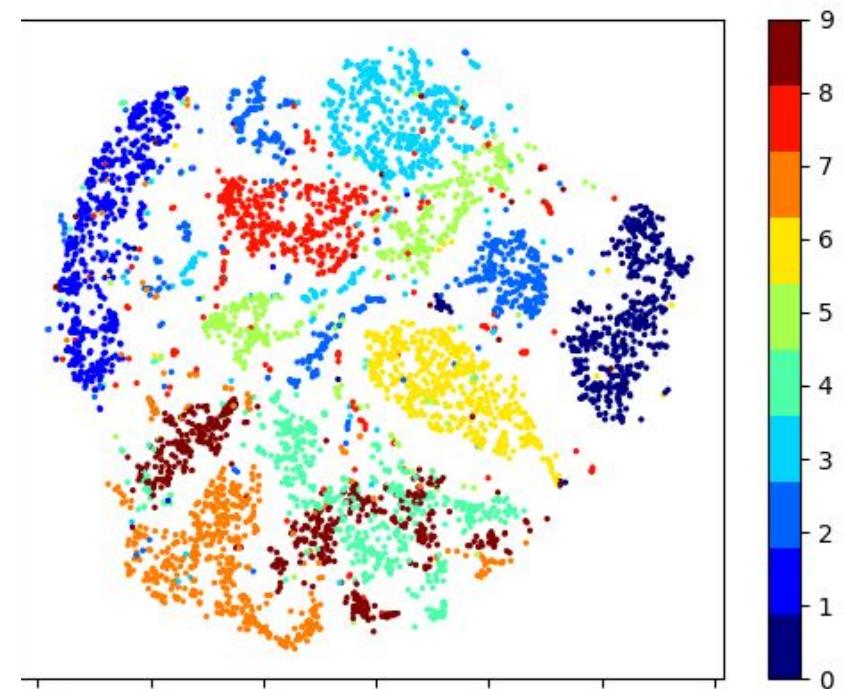


Multi-class Classification

- Handwritten digits recognition

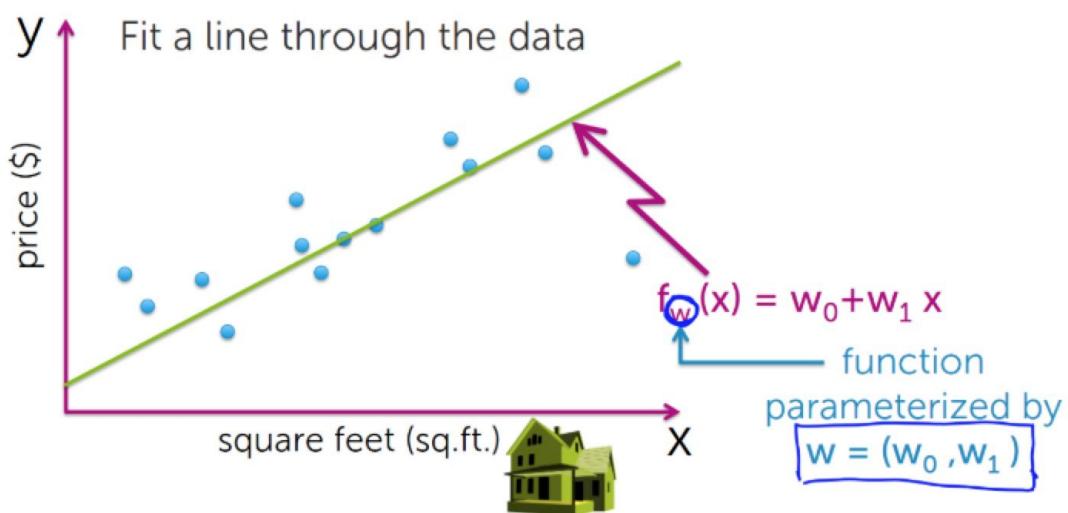


- Ten classes
 - Each number is one class from 0 to 9



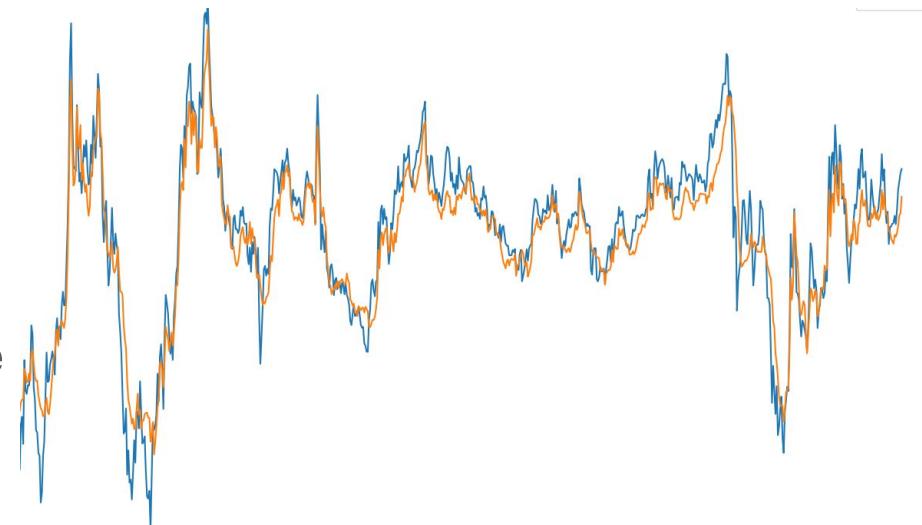
Regression

- Linear Regression



- Non-linear Regression

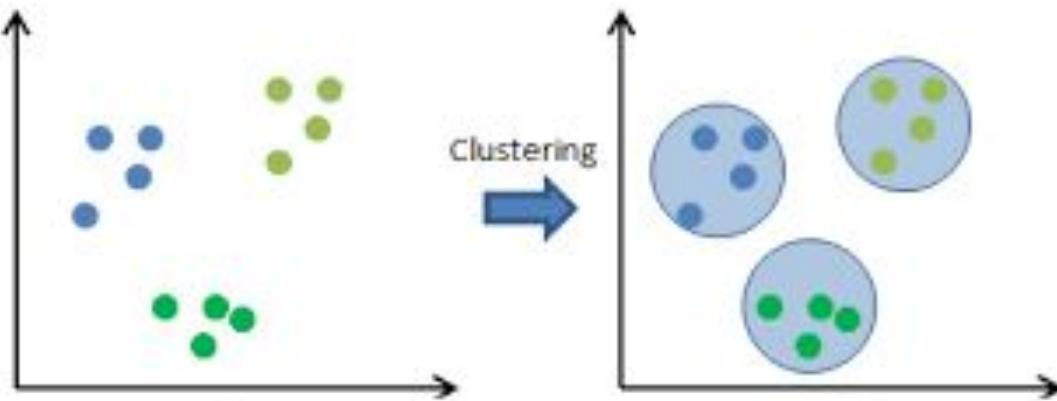
- Time series prediction
- Based on previous k time steps, predict the current value



Unsupervised Learning

- No training labels
- Clustering
 - Divide data into a number of groups
- Dimension Reduction
 - Find a sub feature space for the data representation
- Novelty/Anomaly detection
 - Find the odd one out
- Etc

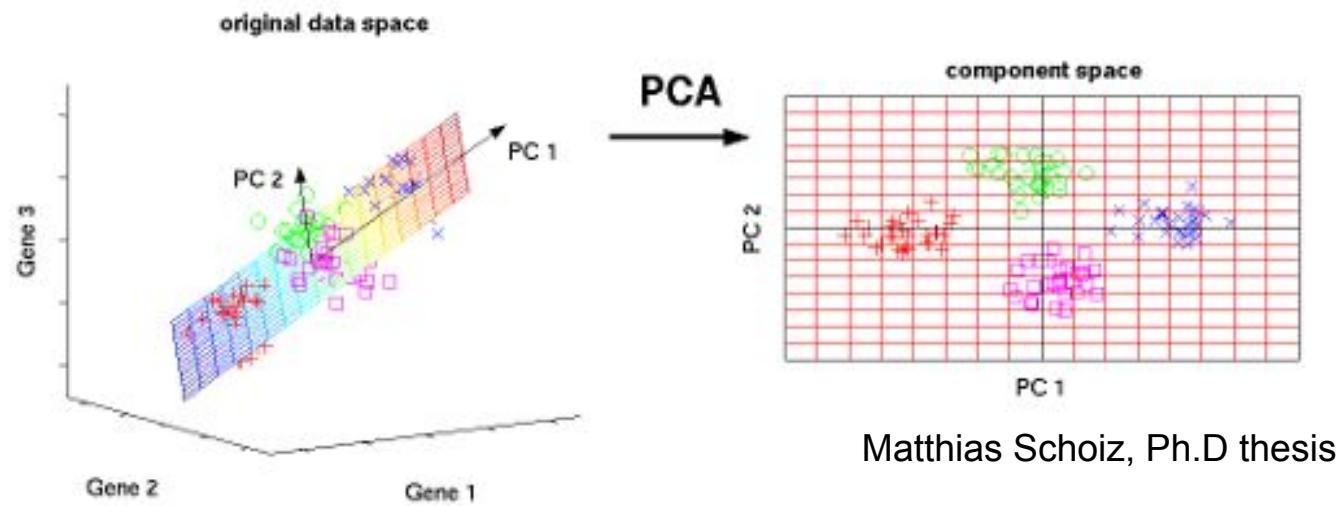
Clustering



- Customer/Marketing segmentation
- Clustering of news, documents, pictures, ..

Dimension Reduction

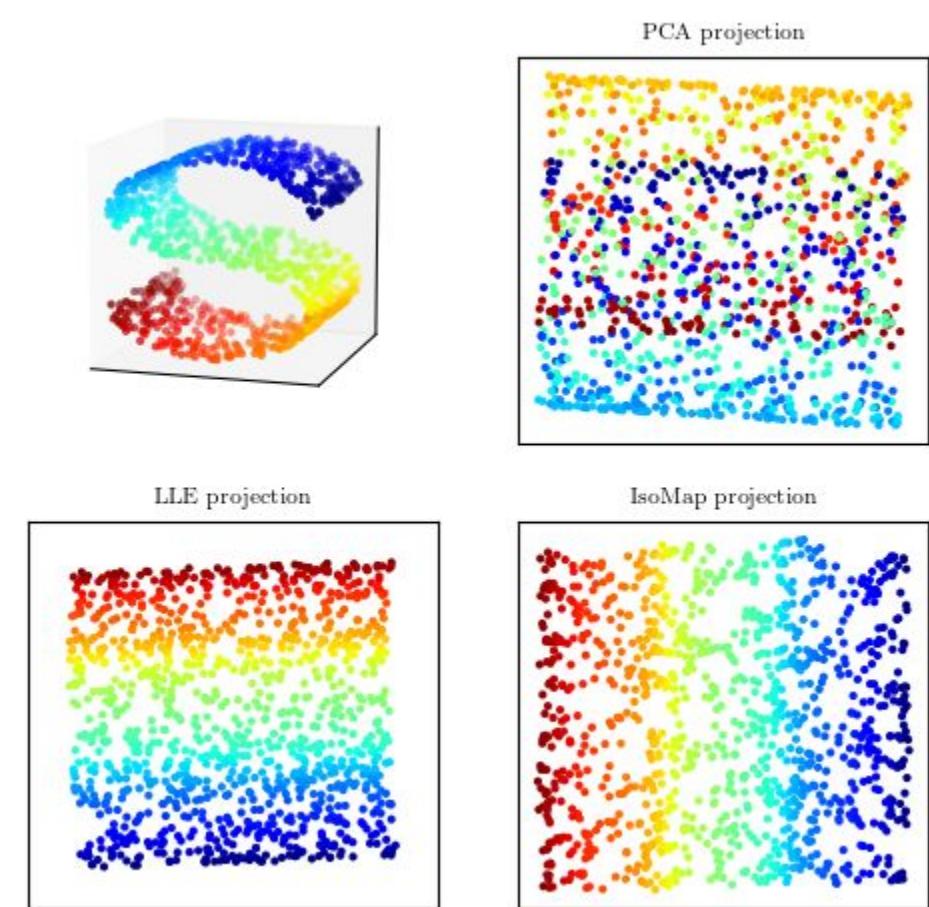
- Principal Component Analysis



Matthias Schoiz, Ph.D thesis

- Manifold Learning

- Find nonlinear subspace/embedding



Novelty / Anomaly Detection

- Identify new/unknown patterns



Interacting with Environment

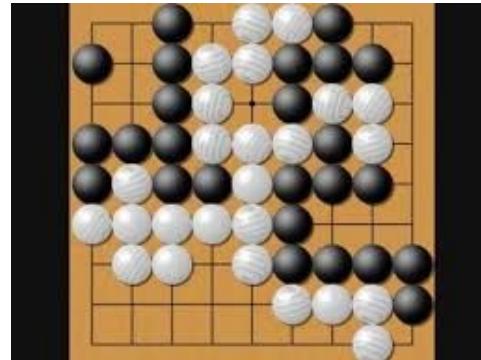
- How to get experience from Environment/How to train the model
- Batch Learning
 - Observe a **batch** of training data $(x_1, y_1), \dots (x_k, y_k)$, then train and used for prediction
- Online Learning
 - **Sequential**: Observe x_1 , predict $f(x_1)$, train with (x_1, y_1) , observe x_2
- Active Learning
 - **interactively** query the user/database to obtain the desired outputs
 - Reduce the amount of labelled data that is needed
- Reinforcement Learning
 - Take action, environment responds, take new actions
 - Play Go, Autonomous Driving

Reinforcement Learning

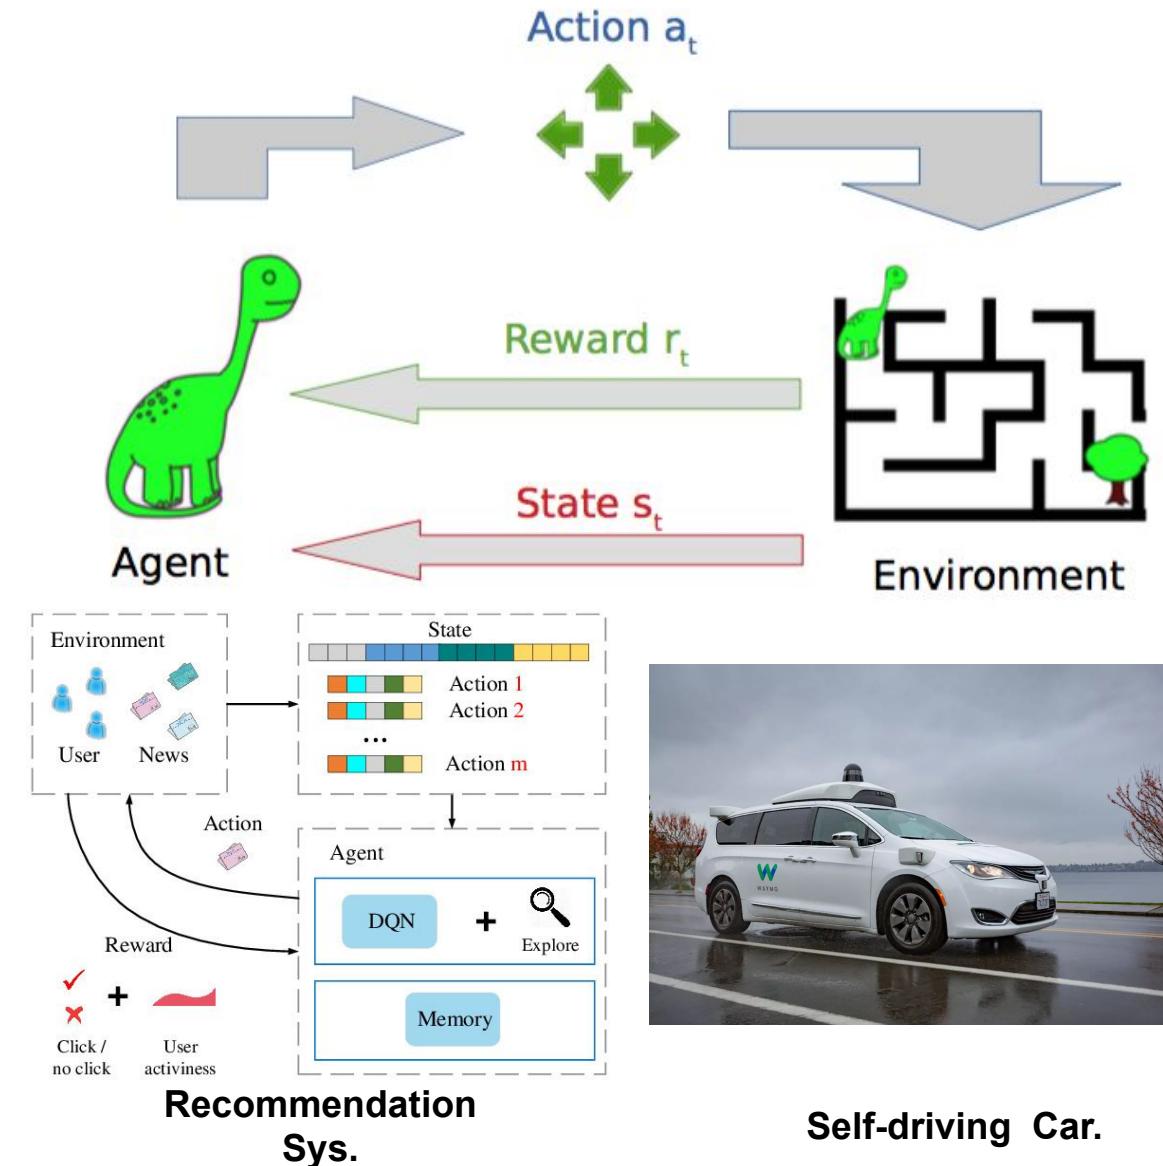
- Do not have known-correct answers
- Repeat
 - Take action
 - Environment reacts
 - Reward/Penalty
 - Update model
- Application



Play Dota2



Play Go



Recommendation Sys.



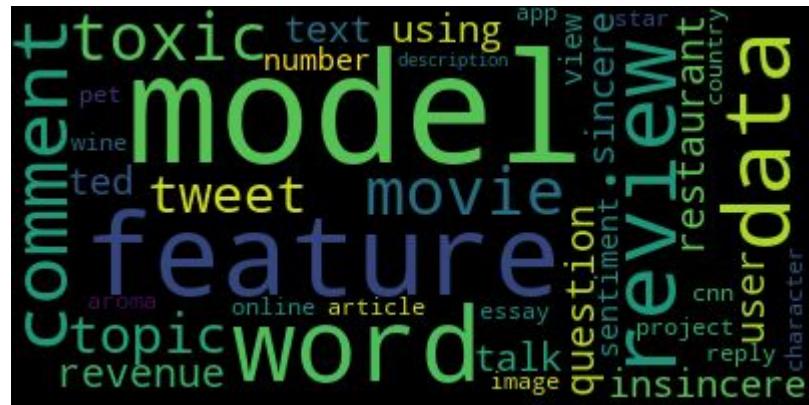
Self-driving Car.

Key Issues in Machine Learning

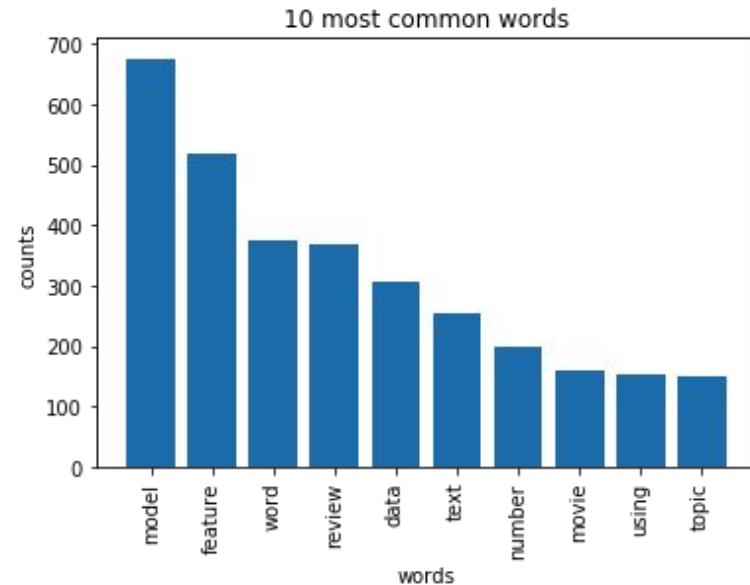
- Obtaining experience
 - How to obtain training data?
 - Supervised or Unsupervised
 - How many examples are enough?
 - PAC learning theory
- Learning algorithms
 - Which kind of machine learning model can approximate function well, when?
 - How does the complexity of learning algorithms impact the learning accuracy?
 - Whether the objective function is learnable?
- Representing inputs
 - How to represent the inputs?
 - How to remove the irrelevant information from the input representation?
 - How to reduce the redundancy of the input representation?

Group Projects

From 2019



Word Cloud of high-frequent words



Topics found via LDA:

Topic #0:

model word feature project tweet movie

Topic #1:

feature model review text data using

Topic #2:

restaurant user review model feature text

Topic #3:

comment feature model character article editor

Topic #4:

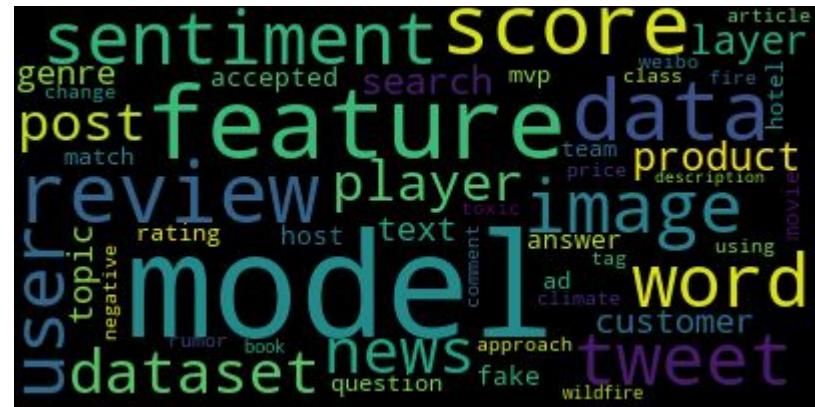
review model movie reply star data

Topic #5:

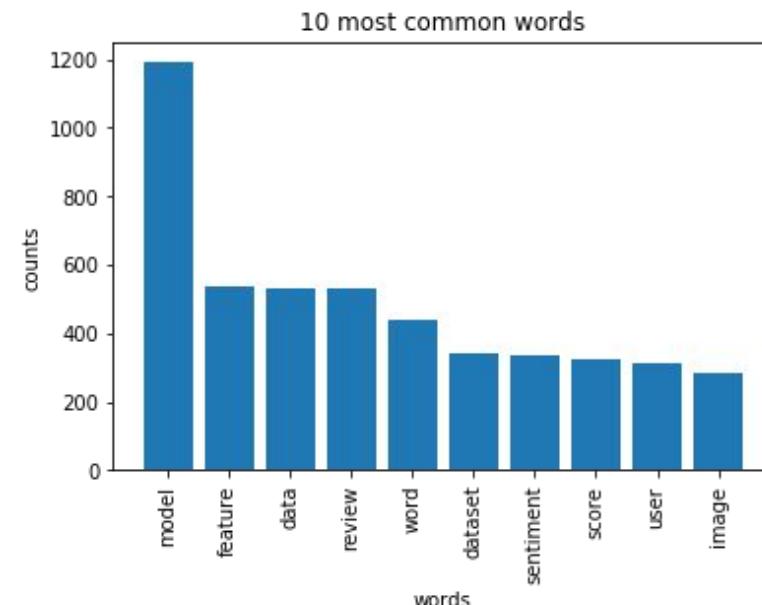
word model data wine tweet question

1. https://bt5153msba.github.io/project/pyp_analyze/2019_reports_Analysis.html
2. <https://bt5153msba.github.io/project/2019fyp.html>

From 2020



Word Cloud of high-frequent words



Topics found via LDA:

Topic #0:
model data word score comment feature

Topic #1:
model data word feature article hotel

Topic #2: image model review dataset product feature

Topic #3:

Topic #4:
review sentiment user negative positive topic

Topic #5:
model answer feature data score sentiment

1. https://bt5153msba.github.io/project/pyp_analyze/2020_reports_Analyses.html
 2. <https://bt5153msba.github.io/project/2020fyp.html>

Project Hint 1

- Find a new problem which can be solved by machine learning technique
 - **Visualize the impact of climate change**



Applying generative models to create personalized images of an extreme climate event, flooding

Source: <https://arxiv.org/pdf/1905.03709.pdf>

Project Hint 2

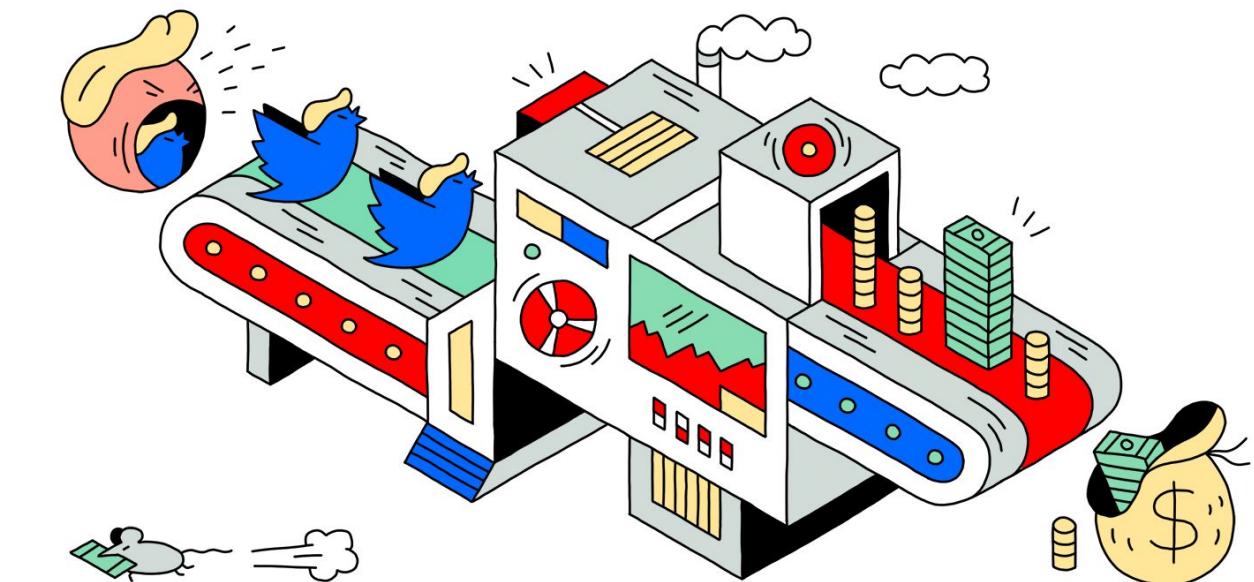
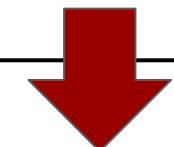
- **Trump2Cash**

1. Monitor Trump's twitter feed

2. Analyze the twitter

If it mentions of any publicly traded stocks
and compute its sentiment

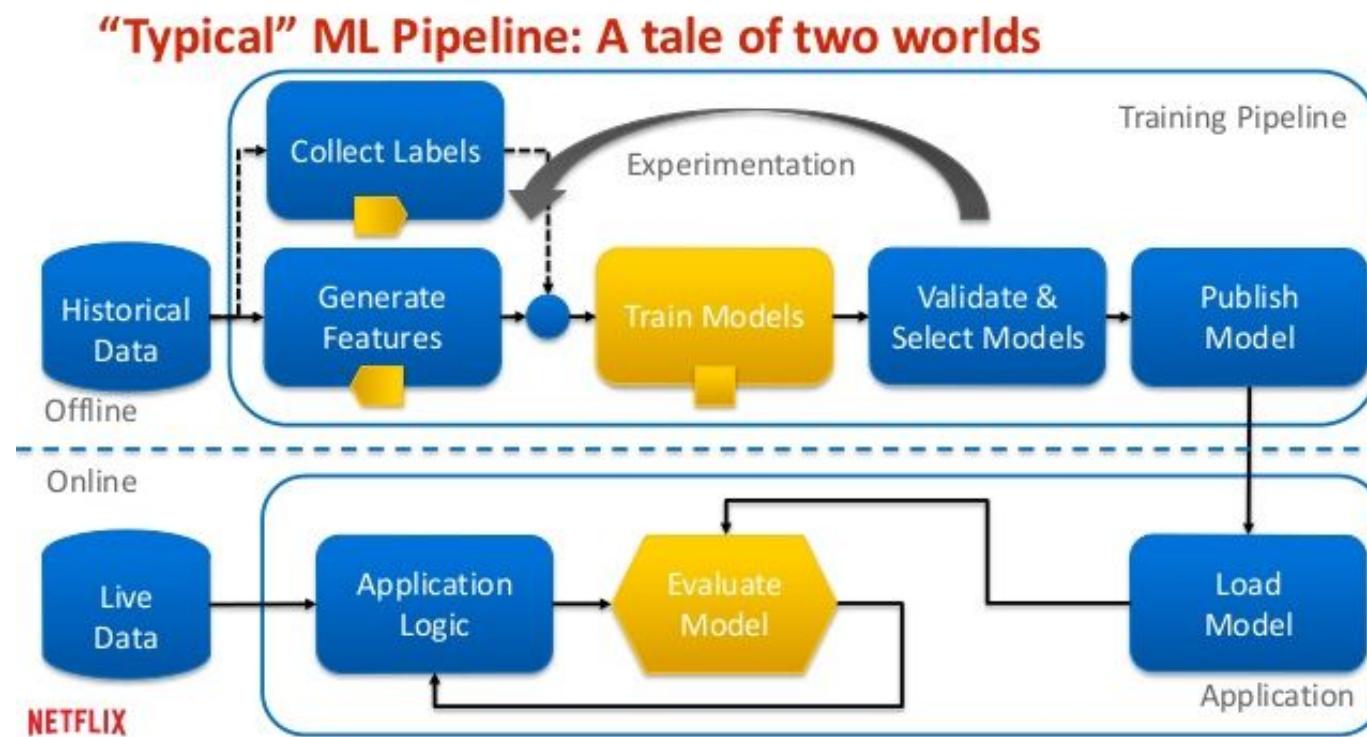
- a. Long it if the sentiment is positive
- b. Short it if the sentiment is negative



Source: <https://github.com/maxbraun/trump2cash>

Project Hint 3

- Build a whole pipeline ML system (real-life one)



Project Hint 4

- In-depth analysis of machine learning algorithms on one specific application
- **Try to explain the findings**

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-rand	76.1	45.0	82.7	89.6	91.2	79.8	83.4
CNN-static	81.0	45.5	86.8	93.0	92.8	84.7	89.6
CNN-non-static	81.5	48.0	87.2	93.4	93.6	84.3	89.5
CNN-multichannel	81.1	47.4	88.1	93.2	92.2	85.0	89.4
RAE (Socher et al., 2011)	77.7	43.2	82.4	—	—	—	86.4
MV-RNN (Socher et al., 2012)	79.0	44.4	82.9	—	—	—	—
RNTN (Socher et al., 2013)	—	45.7	85.4	—	—	—	—
DCNN (Kalchbrenner et al., 2014)	—	48.5	86.8	—	93.0	—	—
Paragraph-Vec (Le and Mikolov, 2014)	—	48.7	87.8	—	—	—	—
CCAE (Hermann and Blunsom, 2013)	77.8	—	—	—	—	—	87.2
Sent-Parser (Dong et al., 2014)	79.5	—	—	—	—	—	86.3
NBSVM (Wang and Manning, 2012)	79.4	—	—	93.2	—	81.8	86.3
MNB (Wang and Manning, 2012)	79.0	—	—	93.6	—	80.0	86.3
G-Dropout (Wang and Manning, 2013)	79.0	—	—	93.4	—	82.1	86.1
F-Dropout (Wang and Manning, 2013)	79.1	—	—	93.6	—	81.9	86.3
Tree-CRF (Nakagawa et al., 2010)	77.3	—	—	—	—	81.4	86.1
CRF-PR (Yang and Cardie, 2014)	—	—	—	—	—	82.7	—
SVM _S (Silva et al., 2011)	—	—	—	—	95.0	—	—

From Yoon Kim