

Applied Machine Learning for Business Analytics

Lecture 2: From BoW to Word2Vec

Agenda

1. Representation Learning in NLP
2. Word Embeddings
3. Neural Networks for NLP
4. Tokens and Embeddings

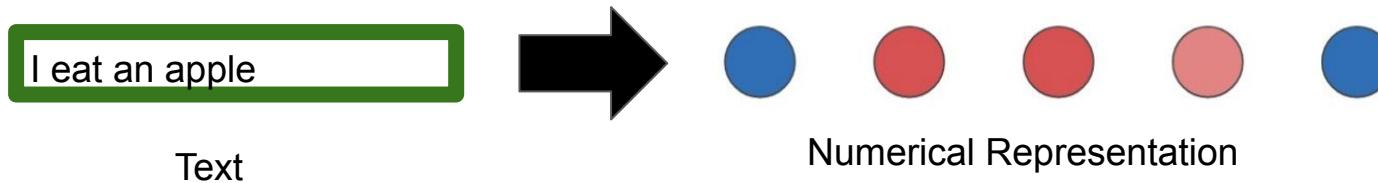
1. Representation Learning

Representation learning

- We need to develop systems that read and understand text the way a person does, by forming a representation of the text, and other context information that humans create to understand a piece of text.

Representation learning

- We need to develop systems that read and understand text the way a person does, by forming a representation of the text, and other context information that humans create to understand a piece of text.



The learned representation should capture high-level semantic and syntactic information.

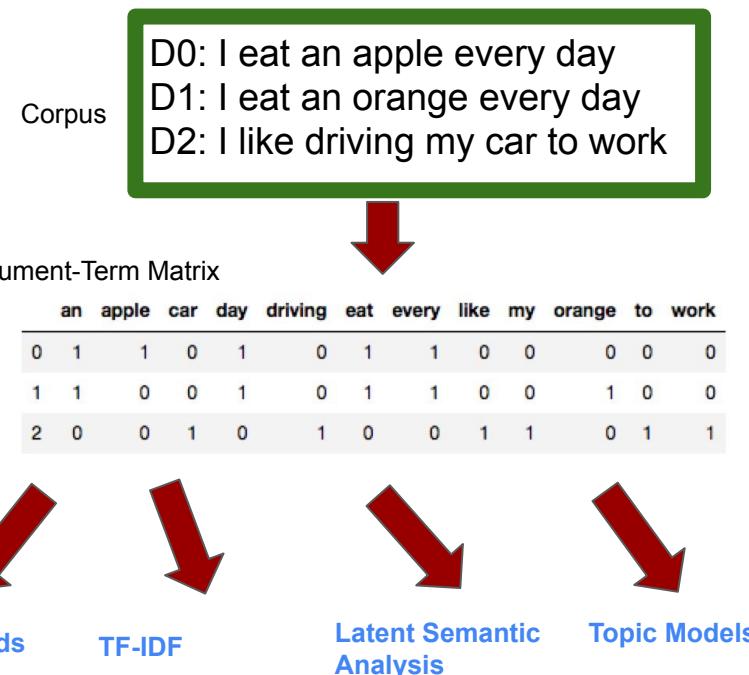
History of NLP

- Now, neural nlp models are able to achieve state-of-arts results in all tasks.
- Before neural nlp:
 - Symbolic NLP: rule-based system (derived from linguistic)
 - Statistical NLP: data-driven and use statistical methods



Statistical NLP

- Starting from Document-Term Matrix
 - It contains the co-occurrence information
 - Bag-of-Words: n-gram as features
 - TF-IDF: frequency of words to measure importance
 - Matrix Decomposition:
 - SVD->Latent Semantic Analysis
 - Probabilistic model-> Topic Model



Bag-of-Words

Building Vocabulary: Tokenization -> Count unique set

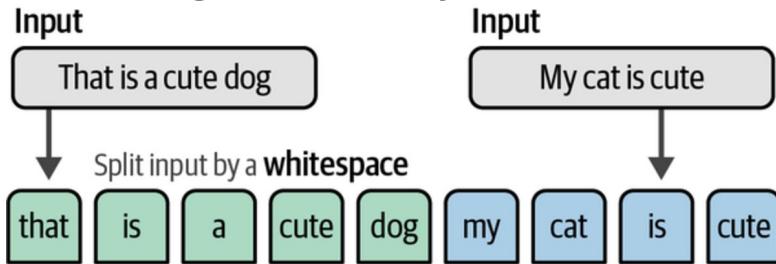


Figure 1-3. Each sentence is split into words (tokens) by splitting on a whitespace.

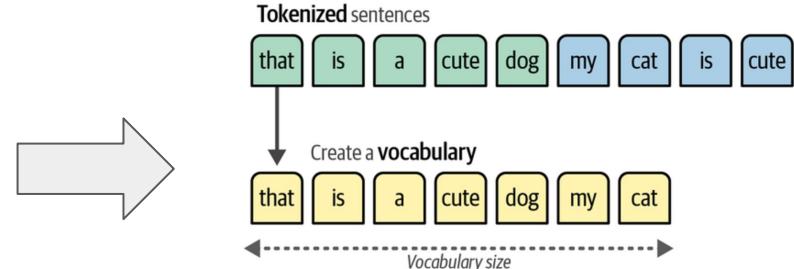


Figure 1-4. A vocabulary is created by retaining all unique words across both sentences.

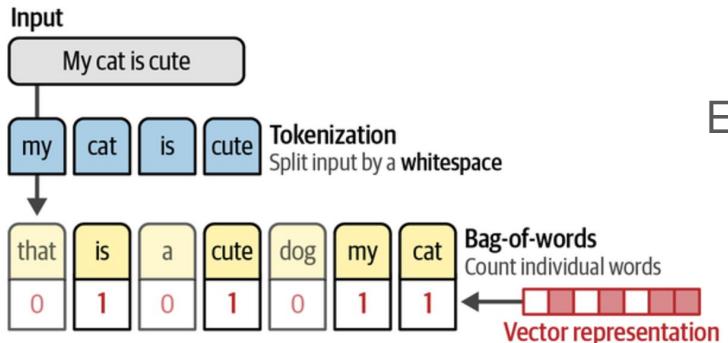


Figure 1-5. A bag-of-words is created by counting individual words. These values are referred to as vector representations.

Encoding Sentences into Vectors

Source: [Hands-On Large Language Models](#)

Limitations of BoW Vectors

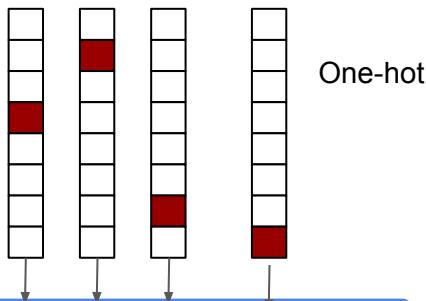
- Too strong assumption: all words are independent of each other
 - $|orange - peach| < |orange - car|$
- Can not capture the order information in the sequence
- High dimensionality due to large size of vocabulary

an	apple	car	day	driving	eat	every	like	my	orange	to	work
0	1	1	0	1	0	1	1	0	0	0	0
1	1	0	0	1	0	1	1	0	0	1	0
2	0	0	1	0	1	0	0	1	1	0	1

A new perspective on BoW

- Each word in vocab is represented in one-hot embedding
- Sum one-hot vectors of the words in a sentence
- The final vector is the representation for the given sentence and then fed into a classifier.

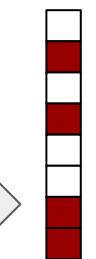
I eat an apple



One-hot

Composition Function:
Sum

BoW
Features



How to improve?

Classification Models

Labels

Statistical NLP

- D3: apple car
 - Word vector: one-hot ones
 - Apple: 0 1 0 0 0 0 0 0 0 0 0 0
 - Car: 0 0 1 0 0 0 0 0 0 0 0 0
 - Sum of two word vectors
 - apple vec + car vec
 - Document vector:
 - 0 1 1 0 0 0 0 0 0 0 0 0

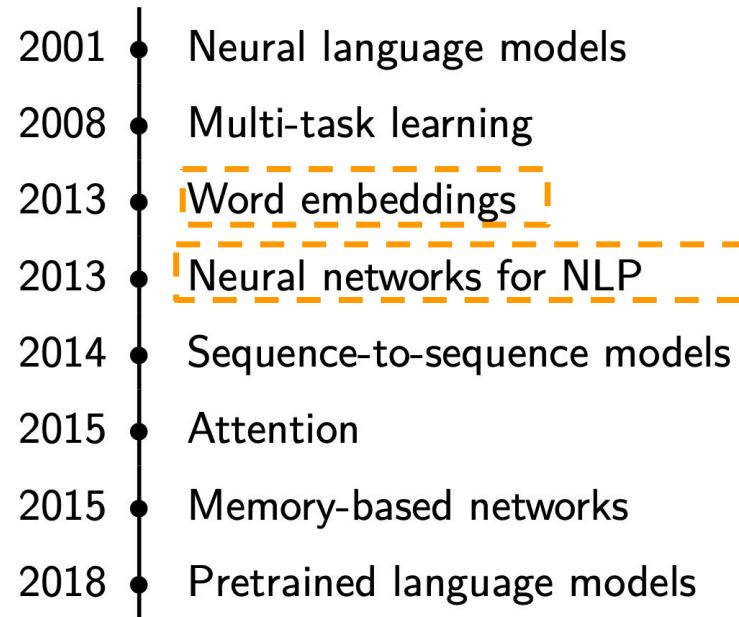
Corpus

D0: I eat an apple every day
D1: I eat an orange every day
D2: I like driving my car to work

Document-Term Matrix

	an	apple	car	day	driving	eat	every	like	my	orange	to	work
0	1	1	0	1	0	1	1	0	0	0	0	0
1	1	0	0	1	0	1	1	0	0	1	0	0
2	0	0	1	0	1	0	0	1	1	0	1	1

Neural NLP



https://www.kamperh.com/slides/ruder+kamper_indaba2018_talk.pdf

2. Word Embeddings

Word representation

- How to represent words in a vector space

apple [0 0 0 0 0 **1** 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0]

orange [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 **1** 0 0 0 0 ... 0 0 0 0 0]

car [0 0 0 0 0 0 **1** 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0]

Distributed representation

- Words should be encoded into a low-dimensional and dense vector

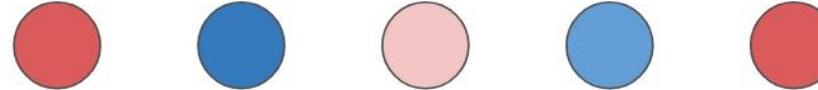
apple



orange



car



Word vectors

Project word vectors in a two-dimensional space. And visualize them!

apple
orange
banana

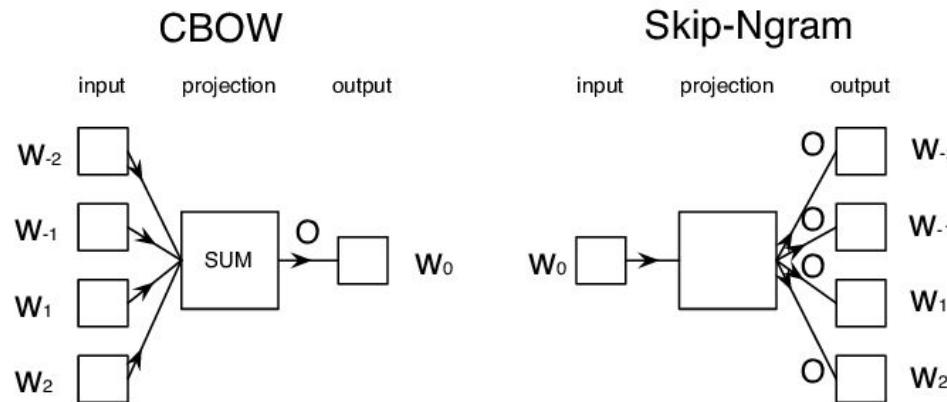
rice
juice
milk

bus
train
car

Similar words are close to each other.

Word2Vec

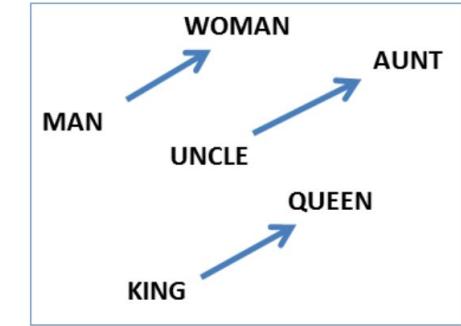
- A method of computing vector representation of words developed by Google.
 - Open-source version of Word2Vec hosted by Google (in C)
 - Train a simple neural network with a single hidden layer to perform word prediction tasks.
 - Two structures proposed Continuous Bag of Words (CBOW) vs Skip-Gram



Word2Vec as BlackBox



input, output



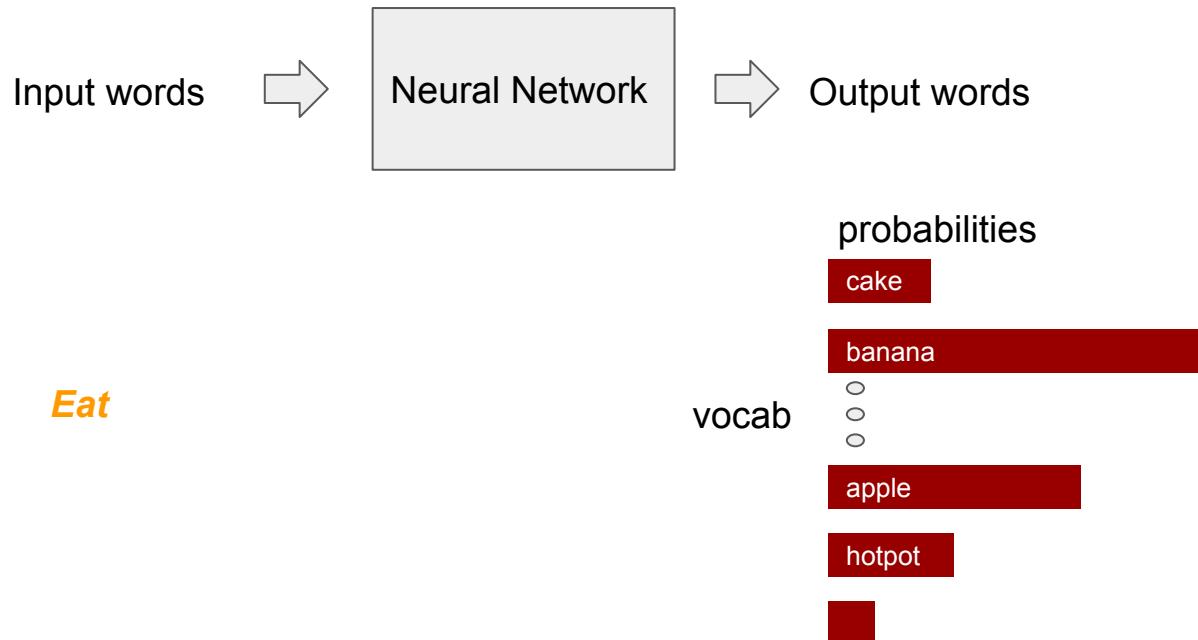
Corpus	Word2Vec Tool	Word Embeddings
A small image of an old document page with vertical columns of Chinese text, representing the input corpus.	A solid black cube, representing the Word2Vec tool as a black box.	A diagram illustrating vector relationships between words. It shows arrows originating from the word "MAN" pointing to "WOMAN", "UNCLE", and "QUEEN". Another arrow originates from "KING" and points to "AUNT". This visualizes how the Word2Vec model captures semantic relationships through vector arithmetic.

Corpus

Word2Vec Tool

Word Embeddings

Use NN to predict word



Self-supervised learning

A Good Visualization for Word2Vec

<https://ronxin.github.io/wevi/>

Target

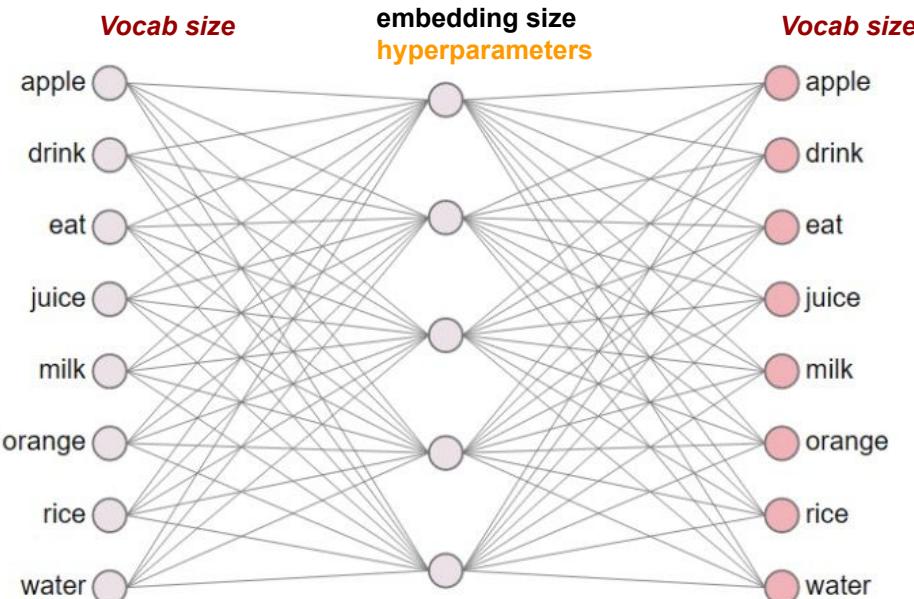
- Given a training corpus, we prepare a list of N (input_word, output_word).
- Objective Function: Maximize probability of all the output words given the corresponding input words.

$$\mathbf{J}(\theta) = \prod_{i=1}^N p(w_{output}^i | w_{input}^i, \theta)$$



**Neural network
parameters that will
be optimized**

Model architecture



Structure Highlights:

- **input layer**
 - one-hot vector
- **hidden layer**
 - linear (identity)
- **output layer**
 - softmax

Input layer

Give the training pair: eat -> apple (given eat, predict apple)

- 8 unique words are in the corpus so that the input layer has 8 neurons
- The index of eat is 3 in the vocab
- The input vector of the $x(\text{eat})$ would be:

One-hot vector

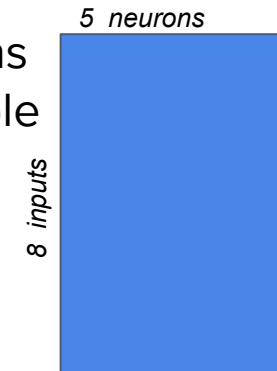
[0,0,**1**,0,0,0,0, 0]

Index of eat

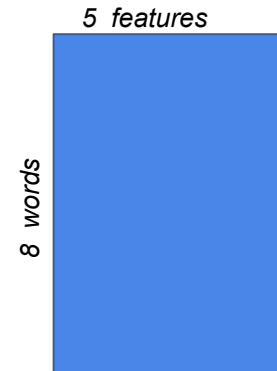
Hidden layer

- **Linear-activation** function here
- **5** neurons are the word vec. dimensions
- This layer is operating as a ‘lookup’ table
- Input word matrix denoted as **IVec**

Hidden Layer Weights Matrix



Word Vector Look Up Table



One-hot vector

[0,0,**1**,0,0,0,0, 0]

Index of eat

1.06	2.91	0.29	1.39	0.33
1.60	1.12	0.29	0.74	0.21
0.96	1.50	1.37	0.34	1.04
0.53	2.11	0.76	2.51	0.20
0.31	0.64	2.08	0.24	1.23
1.40	1.36	0.01	1.69	1.95
2.97	2.13	0.86	0.90	2.21
1.05	0.80	2.18	2.43	1.57



Word vector for “eat”

0.96, 1.5, 1.37, 0.34, 1.04

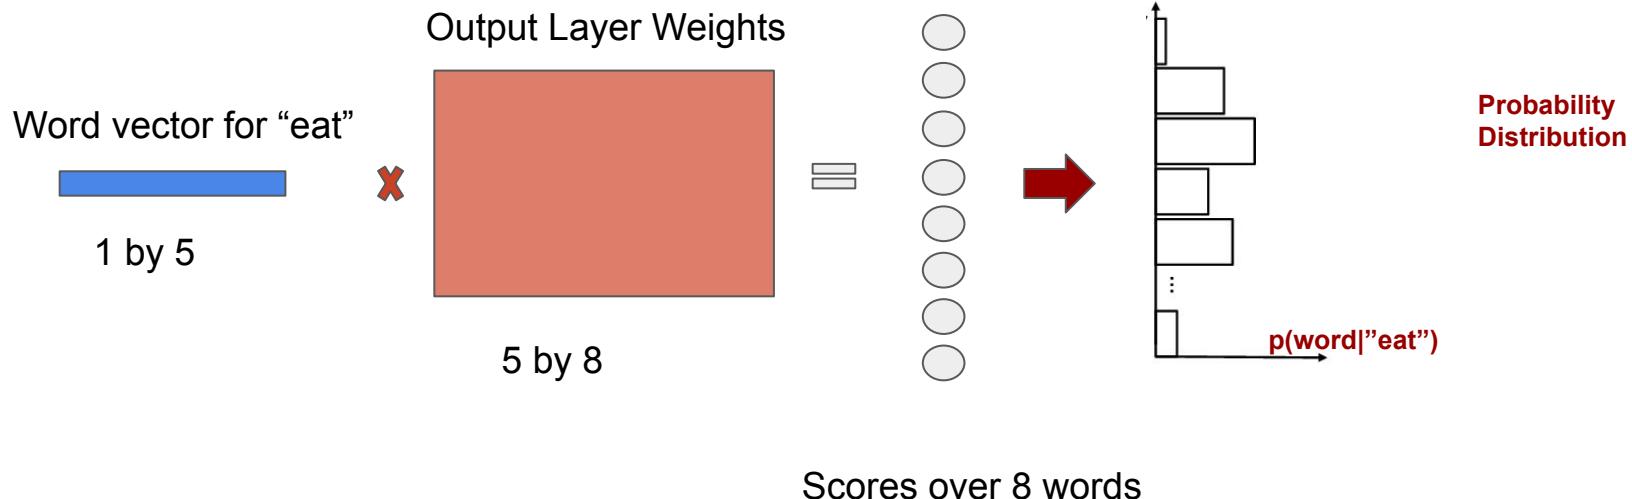
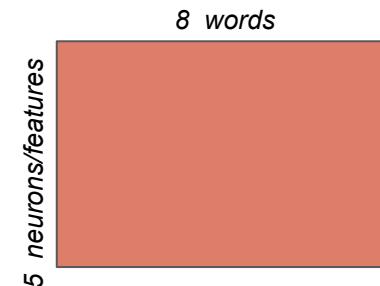


This is a **projection/look up** process: given the index of the word, we take the i th row in the word vector matrix out

Output layer

- Softmax Classifier
- Output word matrix denoted as **OVec**

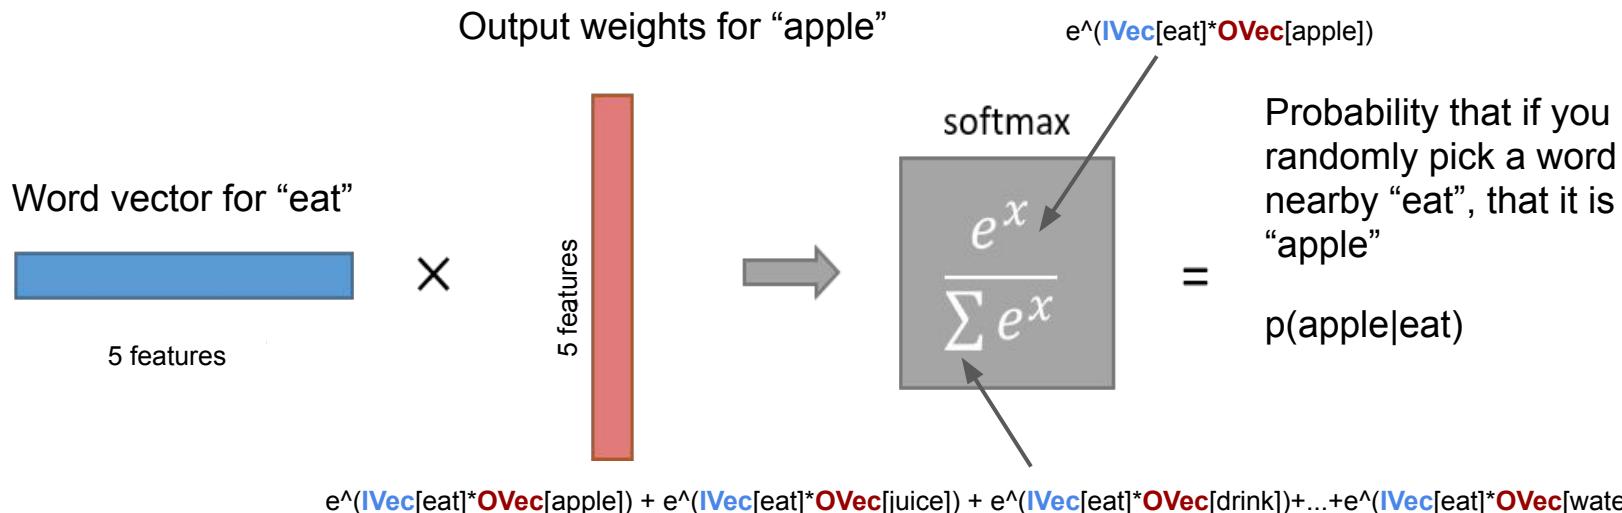
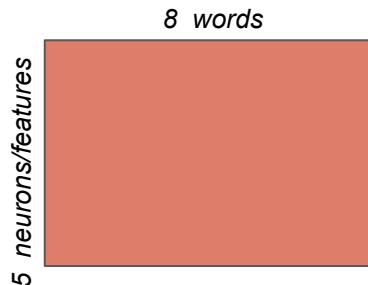
Output Layer Weights Matrix
A.K.A Output word vectors



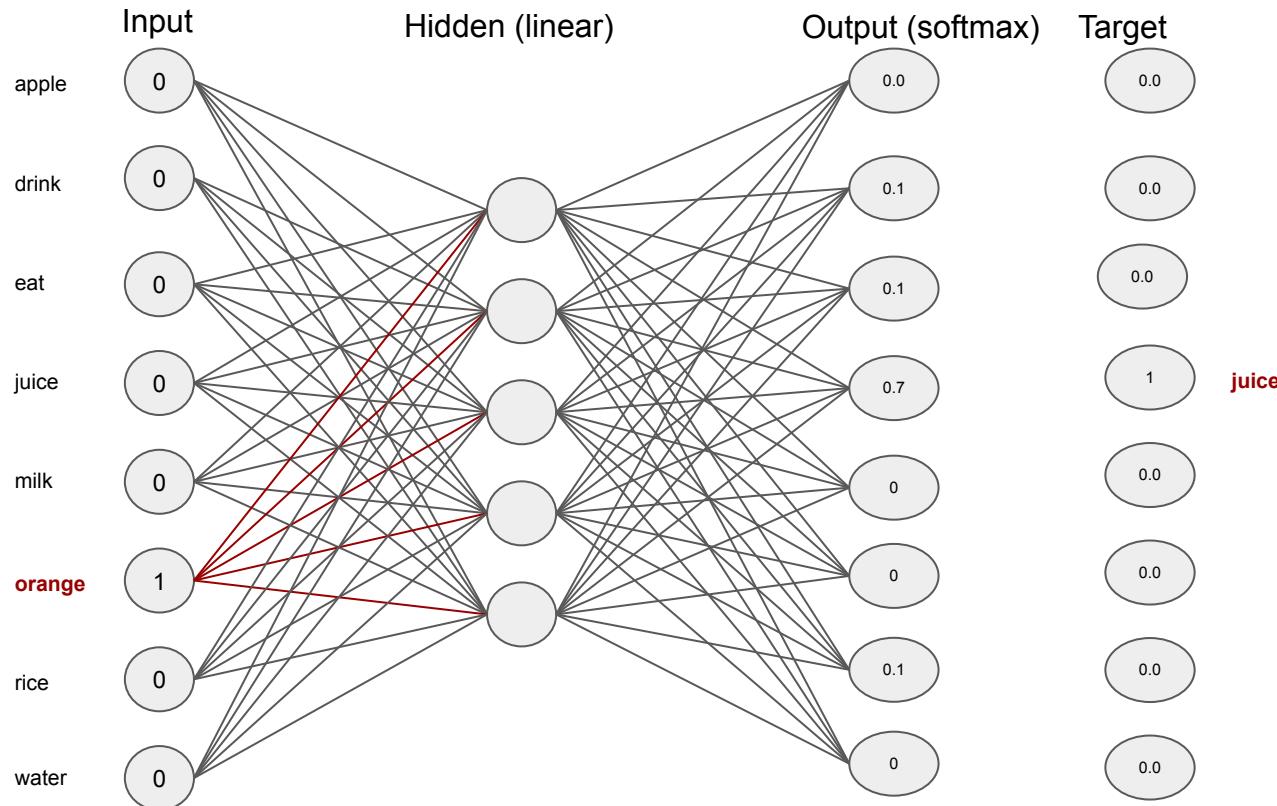
Output layer

- Softmax Classifier
- Output word matrix denoted as **OVec**

Output Layer Weights Matrix
A.K.A Output word vectors

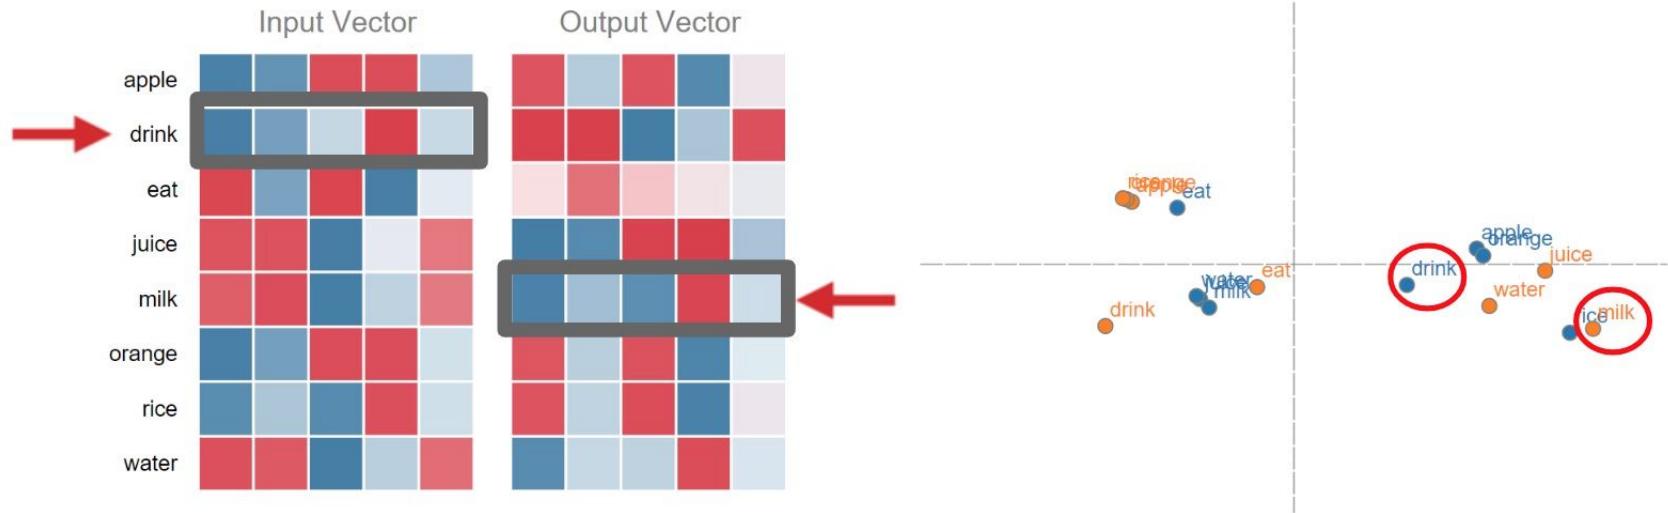


Word2Vec



Then, we can compute the **loss** and call gradient descent to update model parameters.

Updating word vectors



Input vs output word vectors

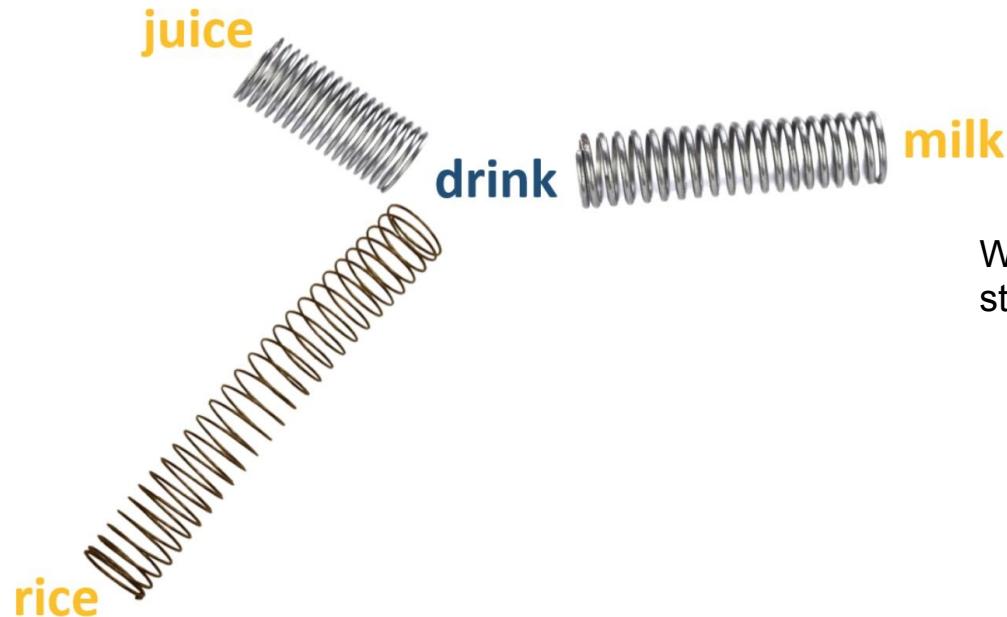
- Input matrix: semantics **encoder** from word index to semantics
- Output matrix: semantics **decoder** from semantics to probability distributions over words
- In most cases, **input** word vectors are used. Some have observed that combinations of these two vectors may perform better

	Vector size	Overall	Semantic	Syntactic
DVRS	300	0.41	0.59	0.26
DVRS	1024	0.43	0.62	0.28
SG	300	0.64	0.69	0.60
SG	1024	0.57	0.60	0.55
Add 300-DVRS, 300-SG	300	0.64	0.72	0.58
Concatenate 300-DVRS, 300-SG	600	0.67	0.74	0.60
Add 1024-DVRS, 1024-SG	1024	0.60	0.66	0.55
Concatenate 1024-DVRS, 1024-SG	2048	0.61	0.68	0.55
Concatenate DVRS-1024, SG-300	1324	0.66	0.73	0.60
Oracle DVRS-1024, SG-300	1024/300	0.70	0.79	0.62

Garten, 2014

Table 2: Performance on word analogy problems with vectors trained against the first 10^9 bytes of Wikipedia.

A force-directed graph



What decides the strength of the string?

Idea behind Word2Vec

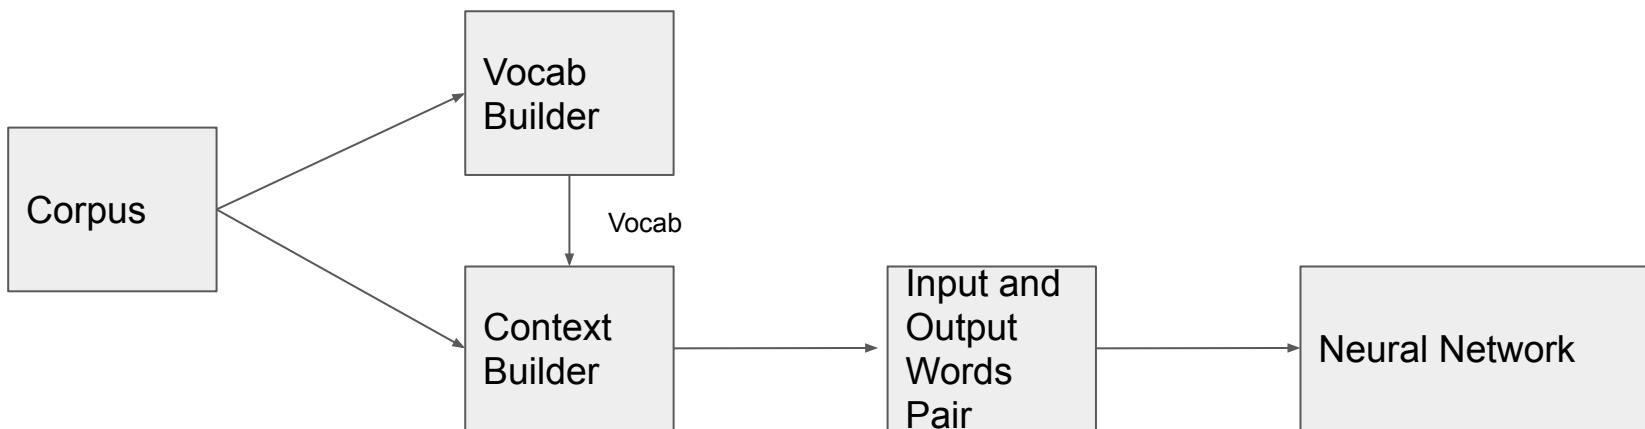
- Feature vector assigned to a word will be adjusted if it can not be used for accurate prediction of that word's context.
- Each word's context in the corpus is the teacher sending error signals back to modify the feature vector.
- It means that words with **similar context** will be assigned **similar vectors!**



“You shall know a word by the company it keeps” - by Firth (1957)

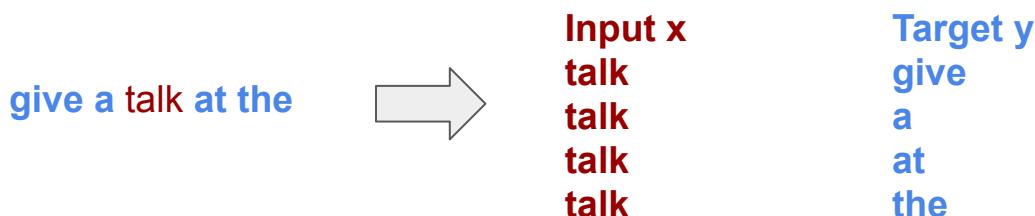
Input and output words

- How to select them from corpus
- Skip-gram and CBoW differ here



Skip-Gram

- Task Definition: given a specific word, predict its nearby word (probability output)
- Model input: source word, Model output: nearby word
- Input is one word, output is one word
- The output can be interpreted as prob. scores, which are regarded as how likely it is that each vocabulary word can be nearby your input word.



CBoW

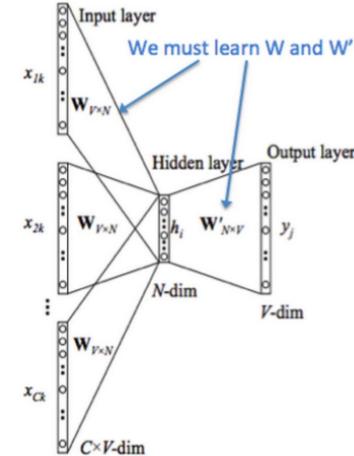
- Task Definition: given context, predict its target word
- Model input: context (several words), Model output: center word
- Input is several words, output is one word
- Core Trick: **average** these context vectors for prob. score computing

give a talk at the



Input x
(give,a,at,the)

Target y
talk

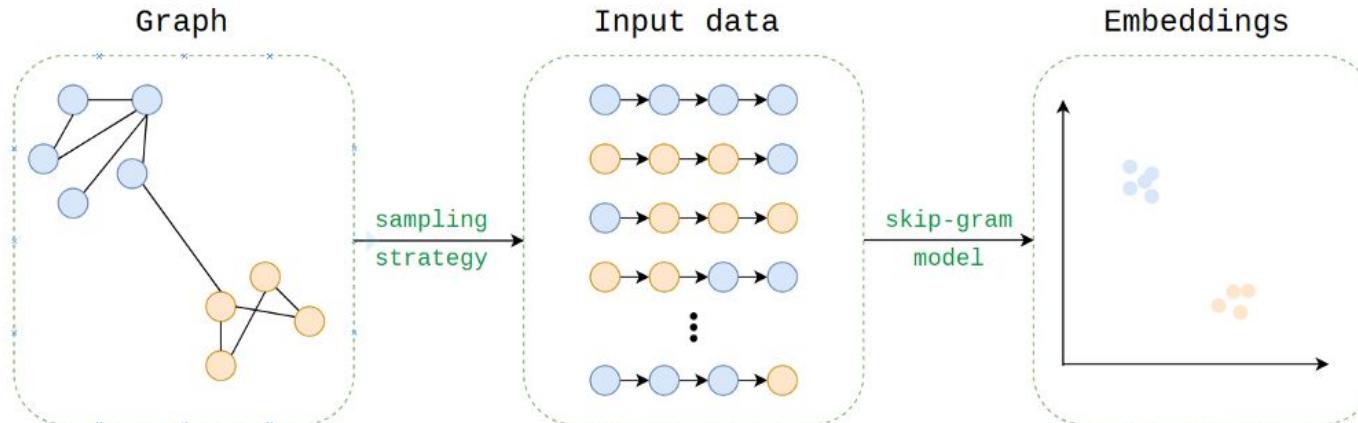


Skip-Gram vs CBoW

- Skip-gram:
 - Learning to predict the context by the center word
 - CBoW:
 - Learning to predict the word by the context
-
- ?: several times faster to train the ?
 - ?: works well with small amount of the training data, represents well even rare words or phrases.

Embedding for graph data

- Embeddings can be extended beyond NLP domain
- Embeddings can be learned for any nodes in a graph
- Nodes can be items, web pages and so on in user clicked stream data
- Embeddings can be learned for any group of discrete and co-occurring states.

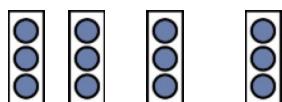


3. Neural Networks for NLP

Sequence of words

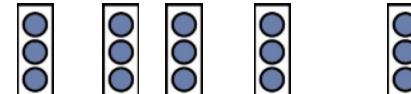
- Each sentence or document can be regarded as a sequence of vectors.
- The shape of matrix depends on the length of sequence. However, the majority of ML systems need fixed-length feature vectors.
- One simple solution: average the sequence of vectors, just like bag-of-words (abandon order information).

I hate this movie



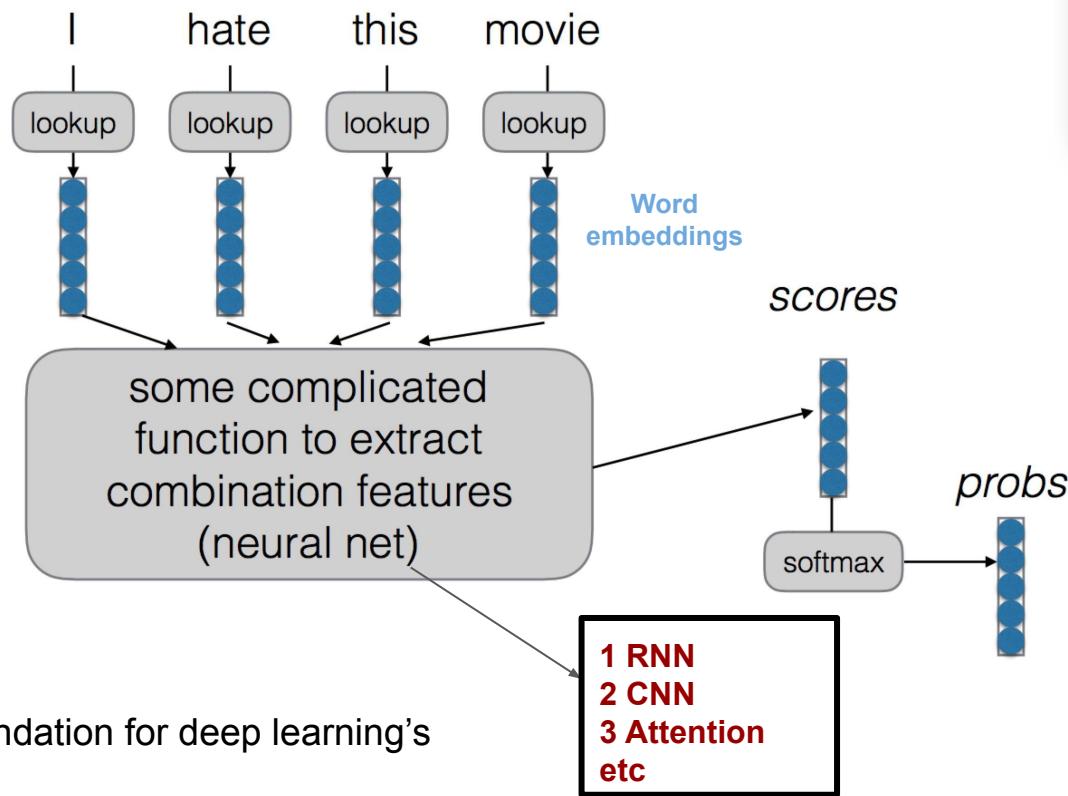
4 by d

This is my favorite movie.

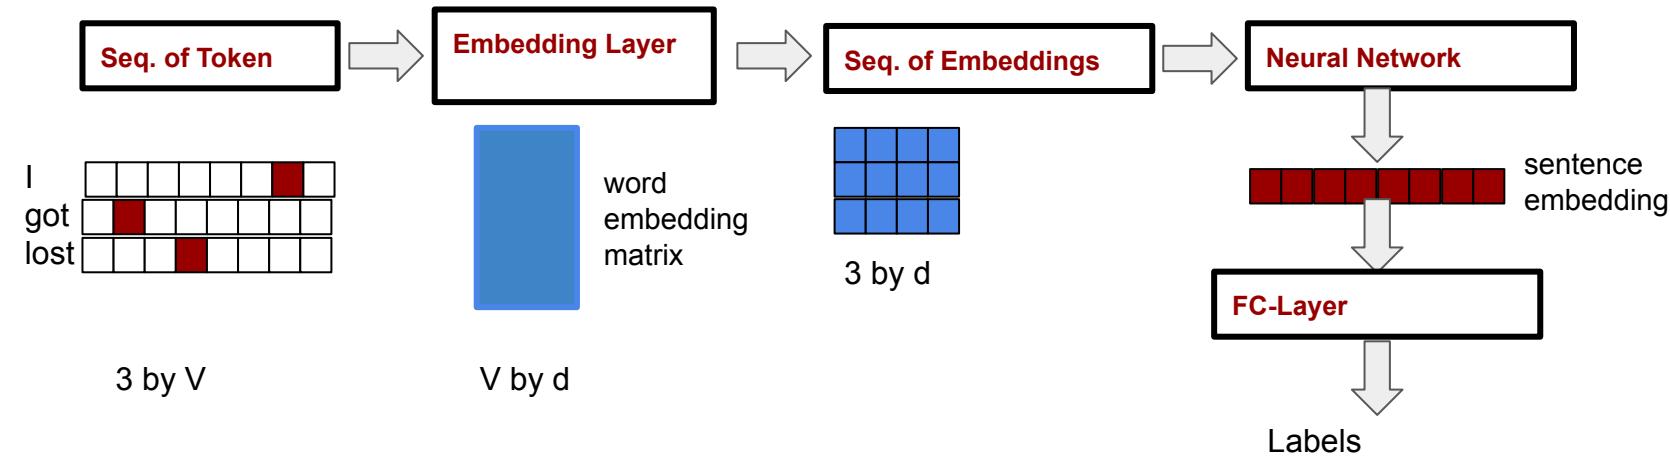


5 by d

Complex semantic



Neural networks for NLP



Embedding layer: the fully-connected layer via one-hot encoding (no bias and no activation)

Is Word2Vec good enough?

- Can not capture different senses of words (context independent)
 - Solution: Take the word order into account->context dependent
- Can not address Out-of-Vocabulary words
 - Solution: Use characters or **subwords**

Multi-sense of Words

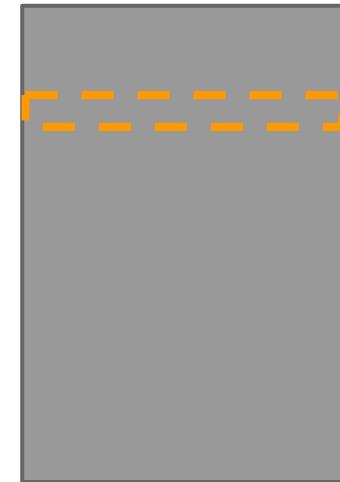
- It is safest to deposit your money in the **bank**.
- All the animals lined up along the river **bank**.
- Today, blood **banks** collect blood.

The third sense of not?

Word2Vec, Fasttext,Glove and other
word embedding models



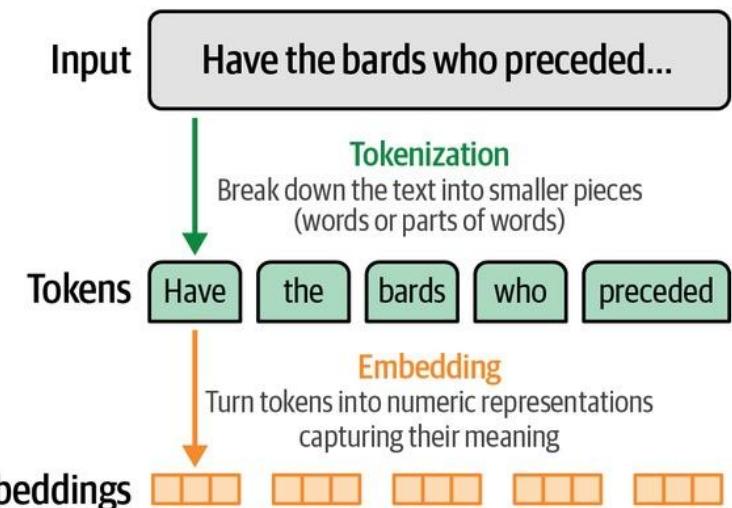
Vocab
size



4. Tokens and Embeddings

Tokens and Embeddings

- Tokenization
 - LLM deal with text in small chunks called tokens.
- Embeddings:
 - The numeric representation for tokens



Tokenization

GPT-4o & GPT-4o mini

GPT-3.5 & GPT-4

GPT-3 (Legacy)

This is business analytics at NUS

Clear

Show example

Tokens

7

Characters

33

This is business analytics at NUS

<https://platform.openai.com/tokenizer>

Text

Token IDs

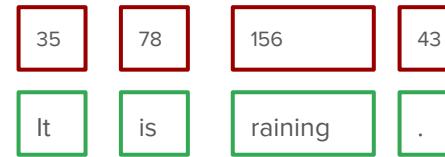
Tokenization Approach

- Word tokens
 - Used in word2vec
 - Unable to deal with new words
 - Result in a vocabulary that has a lot of tokens with minimal differences
 - Apology, Apologize, Apologetic, Apologist
- Subword tokens
 - Contains full and partial words
 - Able to represent new words by breaking down the new token into smaller characters
 - Apolog
 - Suffix tokens: -y, -ize, -etic, -ist

Tokenization: word level

Sentence: It is raining.

Tokens ID



Input tokens

Load tokenizers

Tokenizers are classes that split the sentence into a list of words and subwords defined in the predefined Vocab. The process is given as:

1. Tokenizing: breaking down of the sentence into tokens
2. Adding the [CLS] token at the beginning of the sentence
3. Adding the [SEP] token at the end of the sentence
4. Converting each token into their corresponding IDs in the model

For some words, the BERT tokenization function will first break the word into two subwords, namely characteristic and #Early, where the first token is a more commonly-seen word (prefix) in a corpus, and the second token is prefixed by two hashes ## to indicate that it is a suffix following some other subwords.

```
text = "Hello! How are you doing? I'm at B75153 to learn LLM"

# Basic tokenization
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
print("Original text: ", text)
print("Tokens: ", tokenizer.tokenize(text))
print("Token IDs: ", tokenizer.convert_tokens_to_ids(tokenizer.tokenize(text)))

# Convert to token IDs
tokens_ids = tokenizer.encode(text)
print("Input token IDs: ", tokens_ids)

# Convert back to text
decoded_text = tokenizer.decode(tokens_ids)
print("Decoded text: ", decoded_text)
print("Encoded text: ", [tokenizer.convert_ids_to_tokens(tokens_ids[i]) for i in range(len(tokens_ids))])

# Show special tokens
print("\nSpecial tokens in BERT:")
print("[CLS] token: (ID: 101) (token_id: 101)")
print("[SEP] token: (ID: 102) (token_id: 102)")
print("[PAD] token: (ID: 0) (token_id: 0)")
print("[UNK] token: (ID: 1) (token_id: 1)")

Original text: Hello! How are you doing? I'm at B75153 to learn LLM
Tokens: ['Hello', '!', 'How', 'are', 'you', 'doing', '?', 'I', "'m", 'at', 'B75153', 'to', 'learn', 'LLM']
Token IDs: [101, 955, 999, 229, 2024, 2017, 2725, 1029, 106, 1045, 1085, 2012, 18411, 22243, 22275, 2008, 4953, 2225, 1021, 102]
Decoded text: [CLS] Hello! How are you doing? [UNK] I ' m at B75153 to learn LLM [SEP]
```

Special tokens in BERT:

[CLS] token: (ID: 101)

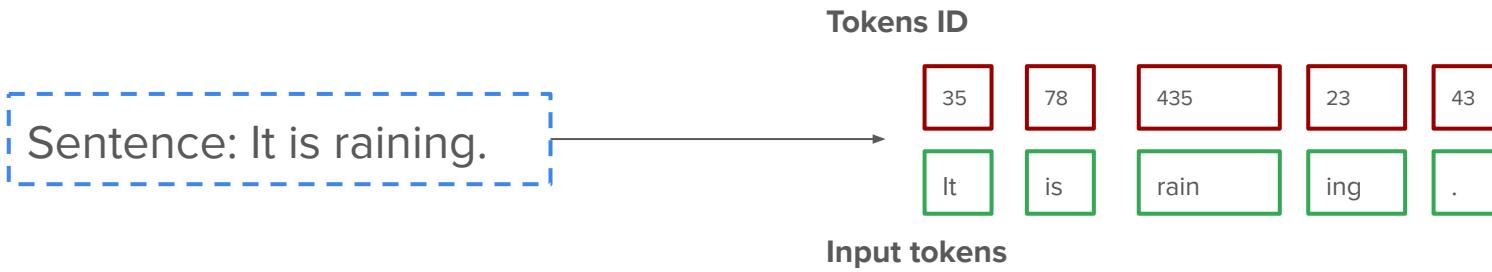
[SEP] token: (ID: 102)

[PAD] token: (ID: 0)

[UNK] token: (ID: 1)

Relies on a predefined vocabulary -> **Out-of-vocabulary issues**

Tokenization: subword level



- Can represent out-of-vocabulary words by composing them from subword units
- The subword algorithms have two main modules:
 - A token learner: this takes a corpus as input and creates a vocabulary containing tokens
 - When should we decompose word into subwords and index those subwords
 - A token segmenter
 - Takes a piece of text and segments it into tokens

Tokenizer Properties

- Tokenization methods
 - How to choose an appropriate set of tokens to represent a dataset
- Tokenizer parameters
 - Vocab size
 - Special tokens
 - Capitalization
- The domain of the data
 - Before the model training, the tokenization method optimized the vocabulary to represent a specific dataset

```
def add_numbers(a, b):
```

.... " " Add the two numbers `a` and `b`. "" "

```
.... return a + b
```

```
def add_numbers(a, b):
```

.... Add the two numbers `a` and `b`.

```
.... return a + b
```

Tokenizer

- Tokenization:
 - Split text into tokens (words, subwords, punctuation, etc.) using **model-specific rules** to match the pretrained model.
- Numerical conversion
 - Convert tokens to number using the **model-specific vocabulary** (indexes), ensuring alignment with the pretrained models

If you do not want to re-train the model, you have to use its associated tokenizers.

```
print("Setting up tokenizer and model...")
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
model = AutoModel.from_pretrained("bert-base-uncased")
```



Bert Tokenizers

Load tokenizers

Tokenize the sentences: split the sentence into a list of word and subwords defined in the predefined Vocab. The process is given as:

1. Tokenization: breaking down of the sentence into tokens
2. Adding the [CLS] token at the beginning of the sentence
3. Adding the [SEP] token at the end of the sentence
4. Converting each token into their corresponding IDs in the model

For some words, the BERT tokenization function will first break the word into two subwords, namely characteristic and ##ally, where the first token is a more commonly-seen word (prefix) in a corpus, and the second token is prefixed by two hashes ## to indicate that it is a suffix following some other subwords.

```
text = "Hello! How are you doing? 🌟 I'm at BT5153 to learn LLM"

# Basic tokenization
tokens = tokenizer.tokenize(text)
print(f"Original text: {text}")
print(f"Tokens: {tokens}")

# Convert to token IDs
token_ids = tokenizer.encode(text)
print(f"\nToken IDs: {token_ids}")

# Convert back to text
decoded_text = tokenizer.decode(token_ids)
print(f"\nDecoded text: {decoded_text}")

# Show special tokens
print("\nSpecial tokens in BERT:")
print(f"[CLS] token: {tokenizer.cls_token} (ID: {tokenizer.cls_token_id})")
print(f"[SEP] token: {tokenizer.sep_token} (ID: {tokenizer.sep_token_id})")
print(f"[PAD] token: {tokenizer.pad_token} (ID: {tokenizer.pad_token_id})")
print(f"[UNK] token: {tokenizer.unk_token} (ID: {tokenizer.unk_token_id})")
```

Original text: Hello! How are you doing? 🌟 I'm at BT5153 to learn LLM
Tokens: ['hello', '!', 'how', 'are', 'you', 'doing', '?', '[UNK]', 'i', "", 'm', 'at', 'bt', '#51', '#53', 't', 'o', 'learn', 'll', '#3m']
Token IDs: [101, 7592, 999, 2129, 2024, 2017, 2725, 1029, 100, 1045, 1005, 1049, 2012, 18411, 22203, 22275, 2000, 4553, 2222, 2213, 102]
Decoded text: [CLS] hello! how are you doing? [UNK] i ' m at bt5153 to learn llm [SEP]

Special tokens in BERT:
[CLS] token: [CLS] (ID: 101)
[SEP] token: [SEP] (ID: 102)
[PAD] token: [PAD] (ID: 0)
[UNK] token: [UNK] (ID: 100)

Next Class: Transformers

Suggested Reading: [The illustrated transformer](#)