

Applied Machine Learning for Business Analytics

Lecture 12: Why do ML Projects Fail in Business

Generative Agents

Generative Agents: Interactive Simulacra of Human Behavior

Joon Sung Park
Stanford University
Stanford, USA
joonspk@stanford.edu

Joseph C. O'Brien
Stanford University
Stanford, USA
jobrien3@stanford.edu

Carrie J. Cai
Google Research
Mountain View, CA, USA
cjcai@google.com

Meredith Ringel Morris
Google Research
Seattle, WA, USA
merrie@google.com

Percy Liang
Stanford University
Stanford, USA
pliang@cs.stanford.edu

Michael S. Bernstein
Stanford University
Stanford, USA
msb@cs.stanford.edu



Figure 1: Generative agents create believable simulacra of human behavior for interactive applications. In this work, we demonstrate generative agents by populating a sandbox environment, reminiscent of The Sims, with twenty-five agents. Users can observe and intervene as agents they plan their days, share news, form relationships, and coordinate group activities.

http://reverie.herokuapp.com/arXiv_Demo/

AutoGPT

- AutoGPT can use GPT4 and GPT3.5 through API to create full-fledged projects by iterating on its own prompts and building upon them in each iteration.
 - Can browse the web
 - Can review the results of its prompts
 - Can create prompts
 - Can read and write files
 - Can install packages
 - Can write code
 -

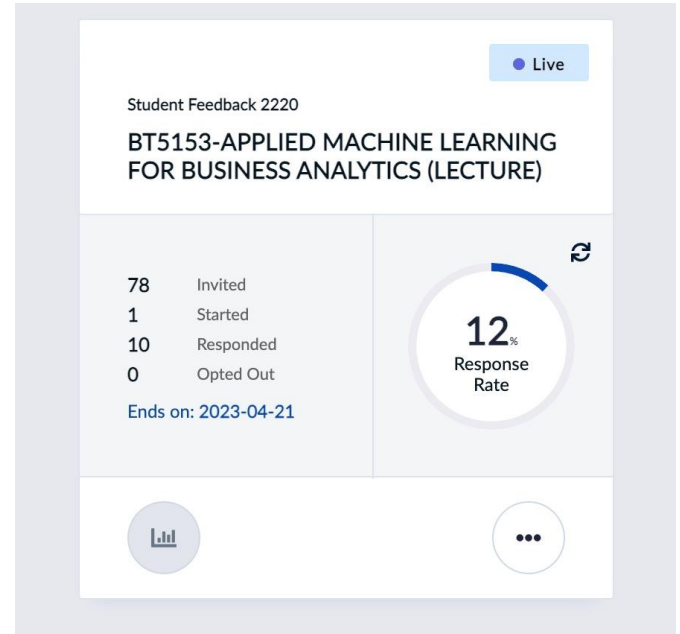
<https://github.com/Torantulino/Auto-GPT>

Action -> Observe outcomes -> Think



Your Voice Matters

- Your Voice Matters
 - BT5153 requires training data for better performance in the future
 - “Student Feedback” under Account in Canvas
 - Thanks ahead for your time!



Logistics

- Thanks for your participation in the Kaggle Competition
- We have three quizzes for attendance check (2 points for each)
 - Even if you miss all, you will get 4 points for attendance
 - Scores would be released before next Friday
- Group project would be due @ next Sunday, April 23
 - Naming of files!
 - Word and latex templates have been provided

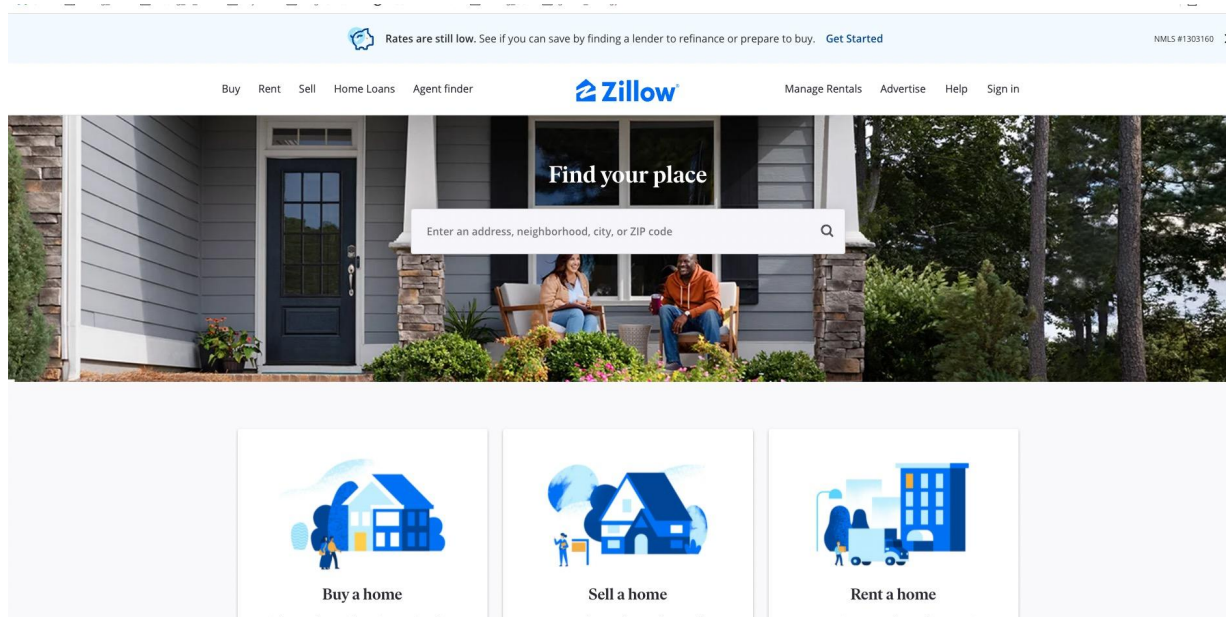
Agenda

1. Zillow Offers
2. Causes of ML Failures
3. Understanding Data Science and Analytics Roles
4. Courser Summary

1. Zillow Offers

Zillow

- an America online real-estate marketplace company



Zillow's Business Model



- How does Zillow make money?
 - <https://seoaves.com/zillow-business-model-how-does-zillow-make-money>
- Zillow Offers:
 - The system called “Zestimate” will analyze multiple data and predict the housing price as the bids for the seller
 - A home sale can be completed in a matter of hours
 - Due to the speed and convenience of the sale process, zillow can purchase houses below market value. After repairs and simple renovation, the houses could be sought to a new buyer at a higher price.
 - The delta in the price + commission fees charged from buyer and seller are the gross profits

1M Kaggle Competition

- At 2018, Zestimate launched a one-million kaggle competition
- In the competition, you will build machine learning models to predict the log error between the actual sale price and the Zestimate based on all the features of a home

$$\text{logerror} = \log(\text{Zestimate}) - \log(\text{SalePrice})$$

The screenshot shows the 'Zillow Prize: Zillow's Home Value Prediction (Zestimate)' competition page. At the top, it states the prize is \$1,200,000 and asks 'Can you improve the algorithm that changed the world of real estate?'. It mentions 3,770 teams and that the competition started 4 years ago. The page has tabs for Overview, Data, Code, Discussion, Leaderboard, and Rules. A 'Late Submission' button is visible. Below the navigation bar, there's a sidebar with links to Overview, Description, Evaluation, Prizes, Timeline, and Competition Overview. The main content area features a large image of a modern house at 111 Archer Ave, New York, NY 10031. The house has 4 beds, 3 baths, and 3,410 sqft. It was built in 2009. The listing includes the asking price of \$1,175,000, the Zestimate of \$1,279,448, and an estimated mortgage of \$4,461/mo. There is also a map showing the location and a contact section on the right.

<https://www.kaggle.com/c/zillow-prize-1>





1M Kaggle Competition

- The winning solution pushed the Zestimate's current nationwide error rate of 4.5% to below 4%

Public Private

The private leaderboard is calculated with approximately 49% of the test data. This competition has completed. This leaderboard reflects the final standings.

Prize Winners

#	△	Team	Members	Score	Entries	Last	Code
1	▲ 9	Zensemble		 0.07408	253	4Y	
2	▲ 1	Juan Zhai 卷宅		 0.07421	53	4Y	

Meet the 'Zillow Prize' winners who get \$1M and bragging rights for beating the Zestimate

BY KURT SCHLOSSER on January 30, 2019 at 5:00 am

[Share](#) 102 [Tweet](#) [Share](#) [Reddit](#) [Email](#)



Zillow Chief Analytics Officer Stan Humphries, left, presents a check to Nima Shahbazi, a member of the team that won the Zillow Prize. (Zillow Photo)



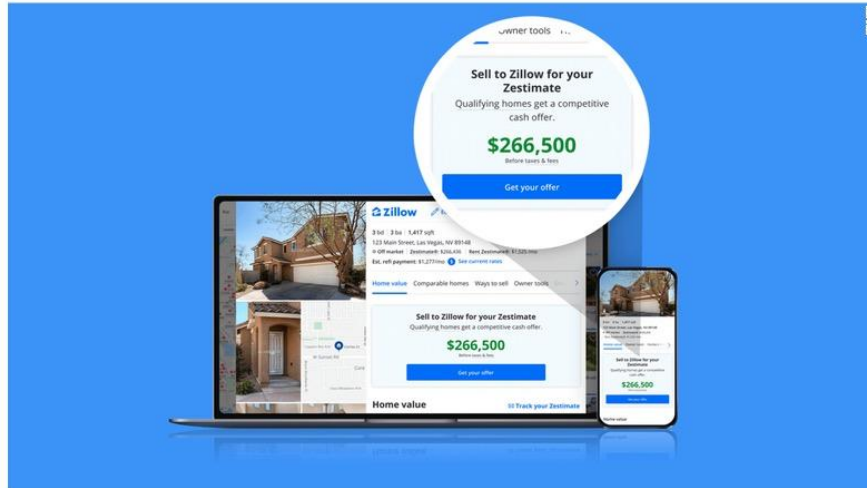
Zestimates was deployed to make offers

Zillow will now make cash offers for homes based on its 'Zestimates'



By Clare Duffy, CNN Business

Updated 1550 GMT (2350 HKT) February 25, 2021



Zillow's "Zestimates" will now represent initial cash offers to homeowners in some markets.

<https://edition.cnn.com/2021/02/25/tech/zillow-zestimate-cash-offer/index.html>

Stock Price from Oct 2021 to Mar 2022

Zillow Group Inc Class C ▲ **54.99** +2.00 (+3.77%)



What went wrong with Zillow Offers?

Zillow, facing big losses, quits flipping houses and will lay off a quarter of its staff.

The real estate website had been relying on its algorithm that estimates home values to buy and resell homes. That part of its business lost about \$420 million in three months.


Zillow is sitting on thousands of houses worth less than what the company paid for them. Caitlin O'Hara for The New York Times

What went wrong with Zillow Offers?

1. Use ML to predict home prices
2. Use predicted prices to flip houses
3. ML models over-predict house prices
4. Buy houses at higher prices

Blaming game

1. Prophet: A python library for forecasting
2. Kaggle-style data science
3. Leadership
4. ML/DS team





Zillow 4.3★

Senior Data Scientist

Washington State

✓ Employer est.: \$127K - \$203K ⓘ

Apply Now

- Proven experience with Forecasting and Time Series modeling, especially Prophet, is strongly preferred.

More insightful discussion:

<https://ryxcommar.com/2021/11/06/zillow-prophet-time-series-and-prices/>

Prophet

- The model is developed by Facebook to predict the web traffic
 - Housing price do not have strong seasonality pattern which is different from web traffic

This is an accurate description of what Prophet's model is. To get a little more in the weeds, Prophet does the following linear decomposition:

- $g(t)$: Logistic or linear growth trend with optional linear splines (linear in the exponent for the logistic growth). The library calls the knots “change points.”
- $s(t)$: Sine and cosine (i.e. Fourier series) for seasonal terms.
- $h(t)$: Gaussian functions (bell curves) for holiday effects (instead of dummies, to make the effect smoother).

Adverse Selection

- Even the model accuracy is high, property owners will only sell when the predicted price is higher than their expected price



2. Causes of ML Failures

ML systems fail silently


Normal softwares fail




Stack Overflow is currently offline for maintenance

Routine maintenance usually takes less than an hour. If this turns into an extended outage, we will tweet updates from [@StackStatus](#) or post details on the [status blog](#).

```
^@^C^A>^D^A^@^P^@^C^AL^D^A^@^T^@^C^A^  
- stack_overflow^M  
^@^C^@R6003^M
```



ML systems fail



Amazon scraps secret AI recruiting tool that showed bias against women

That is because Amazon's computer models were trained to vet applicants by observing patterns in resumes submitted to the company over a 10-year period. Most came from men, a reflection of male dominance across the tech industry.

In effect, Amazon's system taught itself that male candidates were preferable. It penalized resumes that included the word "women's," as in "women's chess club captain." And it downgraded graduates of two all-women's colleges, according to people familiar with the matter. They did not specify the names of the schools.

Japan's Henn na Hotel fires half its robot workforce

“Guests complained their robot room assistants thought snoring sounds were commands and would wake them up repeatedly during the night.”



What is an ML failure?

A failure happens when one or more expectations of the system is violated.

Two types of expectations:

- Operational metrics: e.g. response time, downtime
- ML metrics: e.g. accuracy, MSE, BLUE score (machine translation)

What is an ML failure?

A failure happens when one or more expectations of the system is violated

- Traditional software: mostly operational metrics
- ML systems: operational + ML metrics
 - Ops: returns the risk scores of users within 800ms latency on average
 - ML: Accuracy as 80%

ML system failures

- If you call API to infer the user's risk score and get no response-> ops failure
- If the prediction is incorrect -> ML failure?

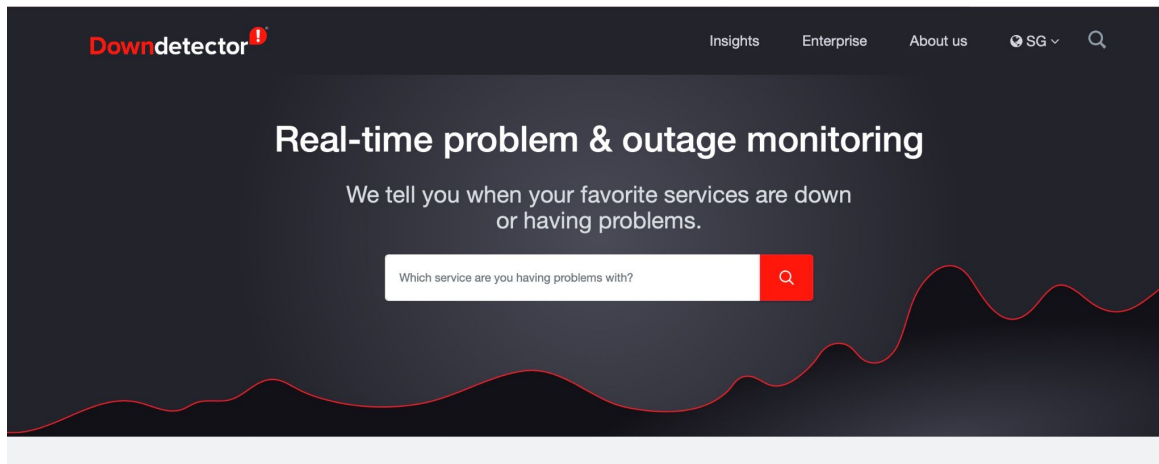
ML system failures

- If you call API to infer the user's risk score and get no response-> ops failure
- If the prediction is incorrect -> Might not be

ML failure when the predictions are **consistently** wrong

What are Ops Failures

- They are normal software systems' failures:
 - Network issues: downtime / crash
 - Deployment issues
 - Hardware issues
 - Dependencies issues



ML-specific failures (during/post deployment)

1. Production data differing from training data
2. Degenerate feedback loops

What are the potential issues for the pre-deployment stage?

Production data differing from training data


- Train-serving skew:
 - Model performing well during development but poorly after production
- Data distribution shifts
 - Model performing well when first deployed, but poorly over time

What is Evidently?

Evidently helps analyze and track data and ML model quality throughout the model lifecycle. You can think of it as an evaluation layer that fits into the existing ML stack.

Evidently has a modular approach with 3 interfaces on top of the shared `analyzer` functionality.


1. Interactive visual reports



<https://github.com/evidentlyai/evidently>

Degenerate feedback loops


- When predictions influence the feedback, which is then used to extract labels to train the next iteration of the model
- Common in tasks with natural labels



**It is not only dangerous
for the business, but
also for society**



Degenerate feedback loops: recsys

- Originally, A is ranked marginally higher than B -> model recommends A
- After a while, A is ranked much higher than B



Degenerate feedback loops: recsys

- Originally, A is ranked marginally higher than B -> model recommends A
- After a while, A is ranked much higher than B



Degenerate feedback loops: resume screening

- Originally, model thinks X is a good feature
- Model only picks resumes with X
- Hiring managers only see resumes with X, so only people with X are hired
- Model confirms that X is good



Replace X with:

- Has a name that is typically used for gender A
- Went to NUS, MSBA

Degenerate feedback loops: resume screening


- Originally, model thinks X is a good feature
- Model only picks resumes with X
- Hiring managers only see resumes with X, so only people with X are hired
- Model confirms that X is good



Tracking feature importance might help!

Detecting degenerate feedback loops

Only arise once models are in production -> hard to detect during training



**Well studied in
Recommendation
System**

Degenerate feedback loops: mitigate

1. Randomization
2. Positional features

Degenerate Feedback Loops in Recommender Systems

Ray Jiang, Silvia Chiappa, Tor Lattimore, András György, Pushmeet Kohli
{rayjiang,csilvia,lattimore,agyorgy,pushmeet}@google.com
DeepMind London, UK

<https://arxiv.org/abs/1902.10730>

3. Understanding Data Science and Analytics Role

Data-Driven Decision

- Understand your business problem
- Identify key challenges and hypothesis in the business problem
- Use data analytics and science methodologies to **test** the hypothesis and **solve** the challenges

How can we make impacts?

- Four stages:
 - How to track/log data?
 - To solve the biz problem, which kind of informations/data are required? -> design trackers
 - It all comes from your business understanding
 - We might need to talk with tech team to design the robust mechanism to make sure the data collection is correct

How can we make impacts?

- Four stages:
 - How to track/log data?
 - How to process data?
 - data cleaning
 - data quality check
 - schema design
 - DE might help to design the raw tables from multiple sources of logs while DSA team needs to design intermedia tables or data mart.

How can we make impacts?

- Four stages:
 - How to track/log data?
 - How to process data?
 - How to analyze data?
 - Dashboard
 - Attribution
 - Metrics design
 - Experimentation
 - Modeling (Data Scientist)

Myth: Modeling is more advanced or DS is on top of DA.

How can we make impacts?

- Four stages:
 - How to track/log data?
 - How to process data?
 - How to analyze data?
 - Dashboard
 - Attribution
 - Metrics design
 - Experimentation
 - Modeling (Data Scientist)

Problem Formulation First, Less Methodology-focused

Simple Tool to solve important problem > Complex tool to solve important problem >> Solve trivial problem

How can we make impacts?

- Four stages:
 - How to track/log data?
 - How to process data?
 - How to analyze data?
 - How to automate the decision-making process?
 - Machine learning is the answer (rule-based system, supervised models and unsupervised models): data -> pattern -> decision

Toy Example

- E-commerce A: CAC in Google Ads per surface enter is 0.01 sgd and the trx conversion rate is 0.1%

Toy Example

- E-commerce A: CAC in Google Ads per surface enter is 0.01 sgd and the trx conversion rate is 0.1%
- Basic Analysis:
 - Marketing cost per transaction: 10 SGD
 - Compare the number across channels

Toy Example

- E-commerce A: CAC in Google Ads per surface enter is 0.01 SGD and the transaction conversion rate is 0.1%
- Advanced Analysis:
 - Cohort analysis: after the first transaction, how many transactions would happen in the following?
 - Product Sense:
 - Another competitor: CAC is 0.1 SGD while the conversion rate is 5%
 - Funnel Analysis
 - A/B Testing
 - User Survey
 - Better Data
 - Users Income
 - Users Demographics Data
 - Other alternative data to support your hypothesis

Advanced analysis is targeted at improving the conversion rate

Open Discussion

- Steve Jobs introduced iphone in 2007
 - It is an art
 - No data support
 - Business context/understanding is the key here
- Classify dog vs Cat
 - It is a science and engineering problem
 - Fit CNN using our training data
 - Deploy the model
- Data Science is art, science and engineering problem



High-impact data science =
Business context (aka art) +
Experimentation/Modeling (aka science) +
Implementation (aka engineering)

High-impact data analysis =

A	B	C	D	E	F
	Observation	Hypothesis	Analysis	Insight	Recommendation
Method	What is a significant data trend you have seen WoW or over the course of the last month/quarter?	What are possible reasons for why this observation is occurring?	How are you going to determine which hypothesis is true or not?	What can we conclude about what is going on after completing the analysis?	What should we do about the insight?

4. Course Summary

Recap: Bridge the Gap

- Introduction to Machine Learning and its **Application**
 - Gap between theory and practice
- Machine Learning **Practices**
 - Prevent potential issues before model deployment
- **Explainable** Machine Learning
 - Interpretability (remember the case: pokemon vs digimon or jpeg vs png)
- Model **Evaluation** in Machine Learning
 - Modeling
- Model **Deployment** in Machine Learning
 - Production
- Why do ML Projects **Fails** in Business
 - Business Understanding/Product Sense First

Recap: DL

- Neural Networks and Deep Learning
- Deep Learning Practices
- Auto-encoders
- Convolutional Neural Networks
- Frontiers in NLP
 - **Self-supervised** Training

Three Steps in Deep Learning

To approximate the true function, define a function **space**

Neural network structures

Learning Representation

Need a **measure** to evaluate the quality of each potential function in the previous space

Evaluate the function






Objective Function

Search the function space to find the best function based on the measure.

Pick the best function

Optimization

Fully-Connected Neural Network



Sigmoid

Hidden Layer 3

Hidden Layer 2

Hidden Layer 1

Input

Recap: Embedding is everything for DL

Text data with Transformers

TheAiEdge.io







Image data with ConvNets



Any data with variational Auto-encoders



Users and products with Recommender engines



source: <https://newsletter.theaiedge.io/>

Recap: Data-driven Decision

- Causal Inference for **Decision** Making
 - Data -> Insights -> Decisions
 - Without decisions, data and Insights are both cheap
 - As a data analyst or data scientist, we need to make Impact, Impact, Impact
 - Identify product-market fit
 - Improve strategies
 - Find directions
 - Fix issues
 - Quantify targets
 - Prioritization

Machine Learning is not ONLY model tuning


CONCLUSIONS

- Model-blind approaches to AI impose intrinsic limitations on the cognitive tasks that they can perform.
- The seven tasks described, exemplify what can be done with models that cannot be done without, regardless how big the data.
- DATA SCIENCE is only as much of a science as it facilitates the interpretation of data -- a two-body problem involving both data and reality.
- DATA SCIENCE lacking a model of reality may be statistics but hardly a science.
- Human-level AI cannot emerge from model-blind learning machines.

16



Judea Pearl



source of the slides:
<http://causality.cs.ucla.edu/blog/wp-content/uploads/2017/12/nips-dec2017-bw.pdf>

Frustrations of the data scientist

Here's why so many data scientists are leaving their jobs

Frustrations of the data scientist!

So to be an effective data scientist in industry it doesn't suffice just to do well in Kaggle competitions and complete some online courses. It (un)fortunately (depending on which way you look at it) involves understanding how hierarchies and politics works in business. Finding a company that is aligned with your critical path should be a key goal when searching for a data science job that will satisfy your needs. However, you may still need to readjust your expectations of what to expect from a data science role.

Source:

<https://towardsdatascience.com/why-so-many-data-scientists-are-leaving-their-jobs-a1f0329d7ea4>

Look at that, it's Fuck This Shit
O'Clock.

Pour me a
vodka anything.



Thank You

- Immense thanks to Xiaohui, Aishik and Chungel!!
- Enjoy having all of **you** in BT5153 this year. Appreciate your hard work!

